

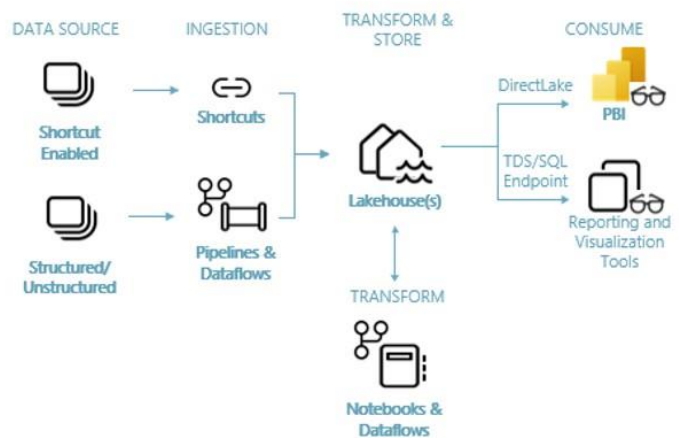
Lakehouse Tutorial

Lakehouse

Published: Jan 2024

Lakehouse

- **Ingestion**
Shortcuts Enabled sources for ADLS G2, Lakehouse, S3, etc. and Pipelines/Dataflows for all other data sources
- **Transformation**
Notebooks and Dataflows
- **Storage**
Lakehouse, or Lakehouse(s) with medallion structure (Bronze/Silver/Gold)
- **Consumption**
In PBI via DirectLake Mode or directly via Lake Warehouse leveraging SQL endpoint
- **Orchestration**
Pipelines triggering Notebook runs and Power BI dataset refreshes



Contents

Introduction..... 3

Module 1: Build your first Lakehouse in Fabric 7

 Create a lakehouse 7

 Data Ingestion..... 9

 Add table to the semantic model..... 18

 Building a report..... 20

Module 2: Ingest, Prep and Analyze 22

 Data Ingestion..... 22

 Data Preparation 30

 Building a report..... 43

Introduction

What is Fabric?

Fabric provides a one-stop shop for all the analytical needs for every enterprise. It covers the complete spectrum of services including data movement, data lake, data engineering, data integration and data science, real time analytics, and business intelligence. With Fabric, there is no need to stitch together different services from multiple vendors. Instead, the customer enjoys an end-to-end, highly integrated, single comprehensive product that is easy to understand, onboard, create and operate. There is no other product on the market that offers the breadth, depth, and level of integration that Fabric offers. Additionally, Microsoft Purview is included by default in every tenant to meet compliance and governance needs.

To get an overview of the components and concepts of Fabric read [Fabric - Overview and Concepts](#).

The purpose of this tutorial

While many concepts in Fabric may be familiar to data and analytics professionals it can be challenging to apply those concepts in a new environment. This tutorial has been designed to walk step-by-step through an end-to-end scenario from data acquisition to data consumption to build a basic understanding of the Fabric UX, the various workloads and their integration points, and the Fabric professional and citizen developer experiences.

The tutorials are not intended to be a reference architecture, an exhaustive list of features and functionality, or a recommendation of specific best practices.

The lakehouse end-to-end scenario

Traditionally, organizations have been building modern data warehouses for their transactional and structured data processing/analytics needs and data lakehouses for big data (semi/unstructured data) processing/analytics needs. These two systems ran in parallel, creating silos, data duplicity and increased total cost of ownership.

Fabric with its unification of data store and standardization on Delta Lake format allows you to eliminate silos, remove data duplicity, and drastically reduce total cost of ownership.

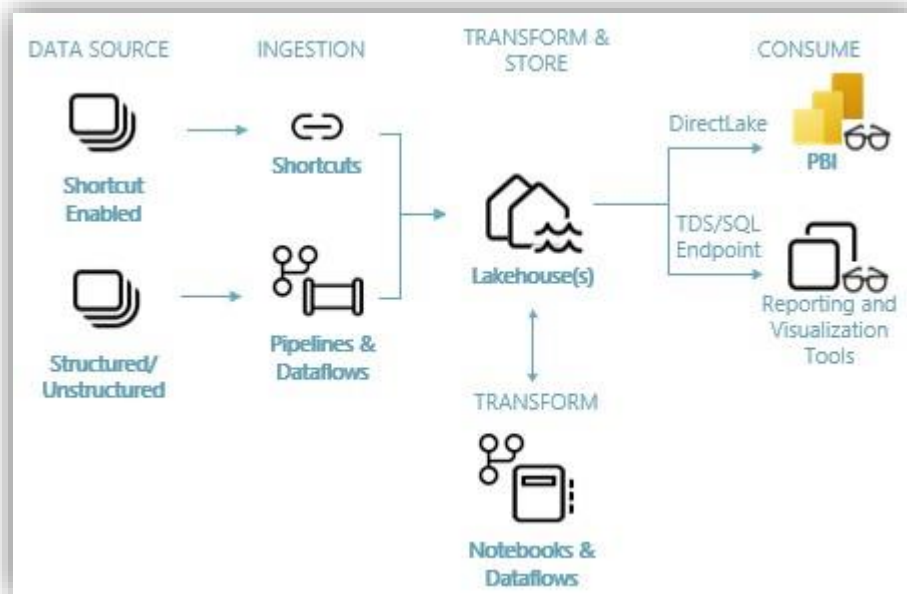
With the flexibility Fabric provides, you can implement either lakehouse or data warehouse architectures or combine these two together to get the best of both the worlds in simple implementation. In this tutorial, you are going to take an example of retail organization and build its lakehouse from start to finish. The same approach can be taken to implement a lakehouse for any organization from any industry.

In this tutorial, you will take on the role of a developer at the fictional Wide World Importers company from retail domain and complete the following steps:

- Sign into your Power BI online account, or if you don't have an account yet, sign up for a free trial.
- Build and implement an end-to-end lakehouse for your organization:
 - Create a Fabric workspace
 - Quickly create a lakehouse – an optional module to implement medallion architecture (Bronze, Silver, and Gold)
 - Ingest, Transform and load data into the lakehouse – bronze, silver and gold zones as delta lake tables for medallion architecture
 - Explore OneLake, OneCopy of your data across lake mode and warehouse mode
 - Connect to your lakehouse using TDS/SQL endpoint
 - Create Power BI report using DirectLake – to analyze sales data across different dimensions
 - Orchestrate and schedule data ingestion and transformation flow with Pipeline

- Cleanup resources by deleting the workspace and other items

The lakehouse end-to-end architecture



Data Sources – Fabric makes it easy and quick to connect to Azure Data Services, other cloud platforms, and on-premises data sources to ingest data from.

Ingestion – With 200+ native connectors as part of the Fabric pipeline and with drag and drop data transformation with dataflow, you can quickly build insights for your organization. Shortcut is a new feature in Fabric that provides a way to connect to existing data without having to copy or move it.

Transform and Store – Fabric standardizes on Delta Lake format, that means all the engines of Fabric can read and work on the same dataset stored in OneLake – no need for data duplicity. This storage allows you to build lakehouses using a medallion architecture or data mesh based on your organizational need. For transformation, you can choose either low-code or no-code experience with pipelines/dataflows or notebook/Spark for a code first experience.

Consume – Data from Lakehouse can be consumed by Power BI, industry leading business intelligence tool, for reporting and visualization. Each Lakehouse comes with a built-in TDS/SQL endpoint for easily connecting to and querying data in the Lakehouse tables from other reporting tools, when needed. When a Lakehouse is created a secondary item, called a Warehouse, will be automatically generated at the same time with the same name as the Lakehouse and this Warehouse item provides you with the TDS/SQL endpoint.

The sample data

For sample data, we are going to use [Wide World Importers \(WWI\) sample database](#). For our lakehouse end-to-end scenario, we have generated sufficient data for a sneak peek into the scale and performance capabilities of the Fabric platform.

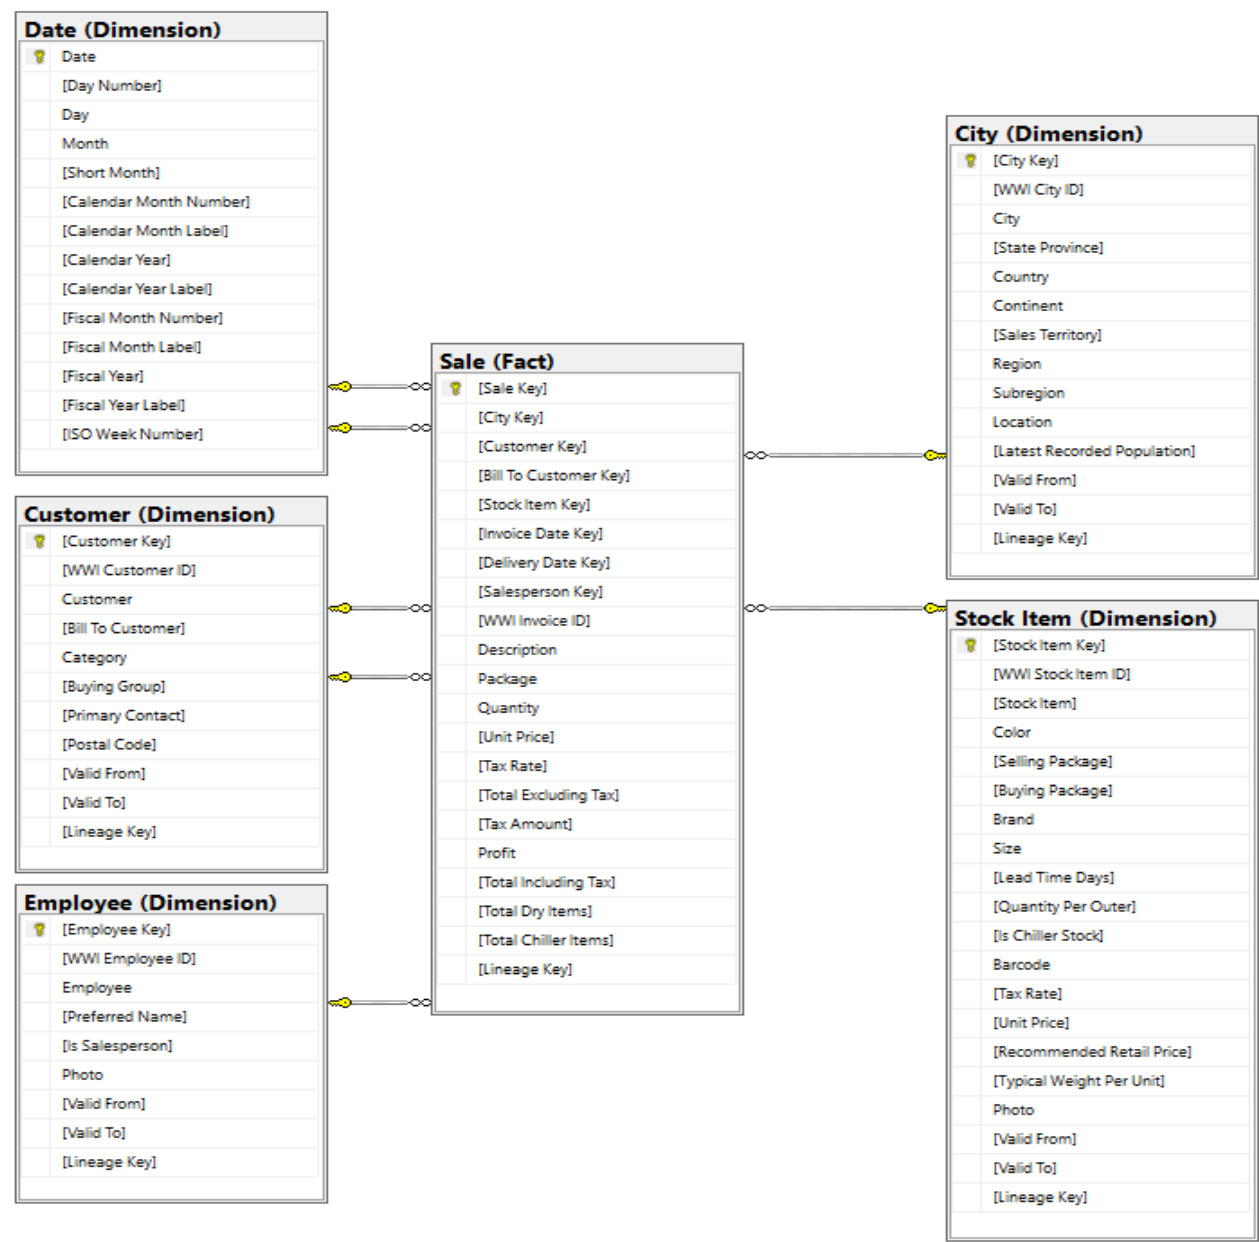
Wide World Importers (WWI) is a wholesale novelty goods importer and distributor operating from the San Francisco Bay area. As a wholesaler, WWI's customers are mostly companies who resell to individuals. WWI sells to retail customers across the United States including specialty stores, supermarkets, computing stores, tourist attraction shops, and some

individuals. WWI also sells to other wholesalers via a network of agents who promote the products on WWI's behalf. You can learn more about their company profile and operation [here](#).

Typically, you would bring data from transactional systems (or line of business applications) into a lakehouse, however for simplicity of this tutorial, we are going to use the dimensional model provided by WWI as our initial data source. We are going to use it as the source to ingest the data into a lakehouse and transform it through different stages (Bronze, Silver, and Gold) of a medallion architecture.

The data model

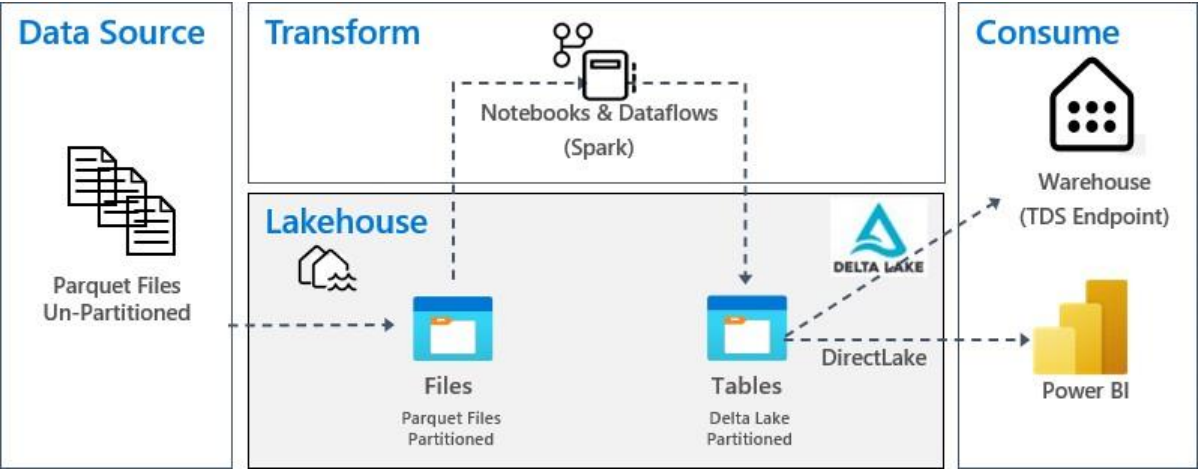
While the WWI dimensional model contains multiple fact tables, for simplicity in explanation we will focus on the Sale Fact table and its related dimensions only, as below, to demonstrate this end-to-end lakehouse scenario:



End-to-end data and transformation flow

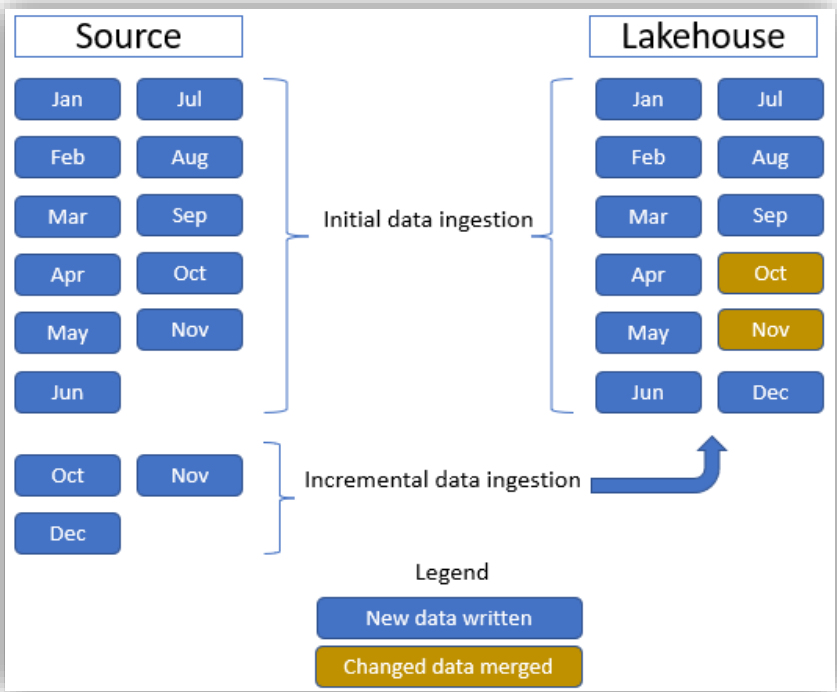
Earlier, we described the [Wide World Importers \(WWI\) sample data](#) which we are going to leverage in building this end to end lakehouse. For this implementation, this sample data is in an Azure Data storage account in Parquet file format for

all the tables, however in your real-world implementation data would likely come from varieties of sources and in a variety of formats.



Data Source – source data is in Parquet file format in an un-partitioned structure, stored in a folder for each table. For the purpose of this tutorial, we will set up pipeline to copy/ingest the complete historical or onetime data to the lakehouse.

To demonstrate process and capabilities for subsequent incremental data load, we will have an optional module at the end of this tutorial. In that module, we use a fact table (Sale) which has one parent folder with all the historical data for 11 months (one subfolder for each month) and another folder for just incremental data for 3 months (one subfolder for each month). During initial data ingestion, 11 months of data is ingested or written into the lakehouse table. However, when the incremental data arrives, it has updated data for Oct and Nov and new data for Dec, so Oct and Nov data are merged and new data for Dec is written into lakehouse table as depicted in figure below.



Lakehouse – For this implementation and for the simplicity’s sake, we are going to create one lakehouse, ingest data into Files section of the lakehouse and then create delta lake tables in the Tables section of the lakehouse.

You will find an optional module at the end of this tutorial which covers creating the lakehouse with medallion architecture (Bronze, Silver, and Gold) and talks about its recommended approach.

Transform – For data preparation and transformation, we are going to demonstrate the use of Notebooks/Spark for code first users as well as demonstrate Pipelines/Dataflow for low-code/no-code users.

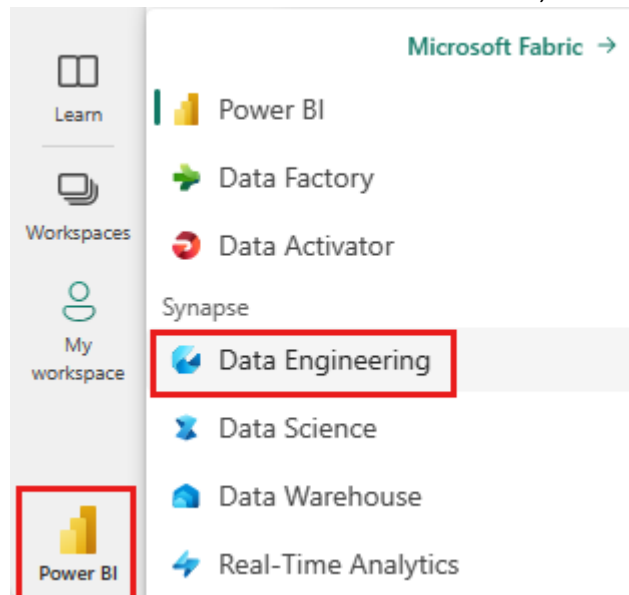
Consume – To demonstrate data consumption you will learn how you can use the DirectLake feature of Power BI to create reports/dashboards and directly query data from the lakehouse. Further, to demonstrate how you can make your data available to third party reporting tools, you can use TDS/SQL endpoint to connect to Warehouse and run SQL-based queries for analytics.

Module 1: Build your first Lakehouse in Fabric

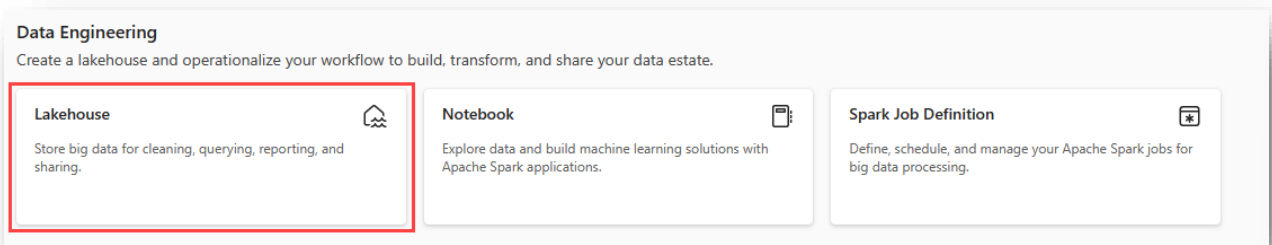
The intent of this module is to quickly build end to end journey of building a lakehouse, ingesting data for a table, applying transformation whenever required and then using ingested data into the lakehouse delta table for creating reports.

Create a lakehouse

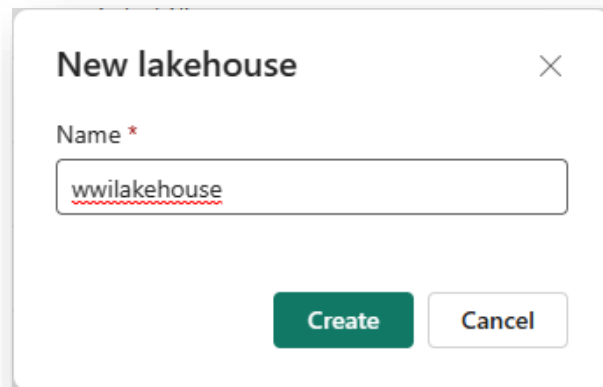
1. In the [Power BI service](#) select **Workspaces** in the left-hand menu.
2. Search for your workspace by typing in the search textbox at the top and click on your workspace to open it.
3. From the workload switcher located at the bottom left of the screen, select **Data engineering**.



4. In the **Data Engineering** section, select **Lakehouse** to create a lakehouse.



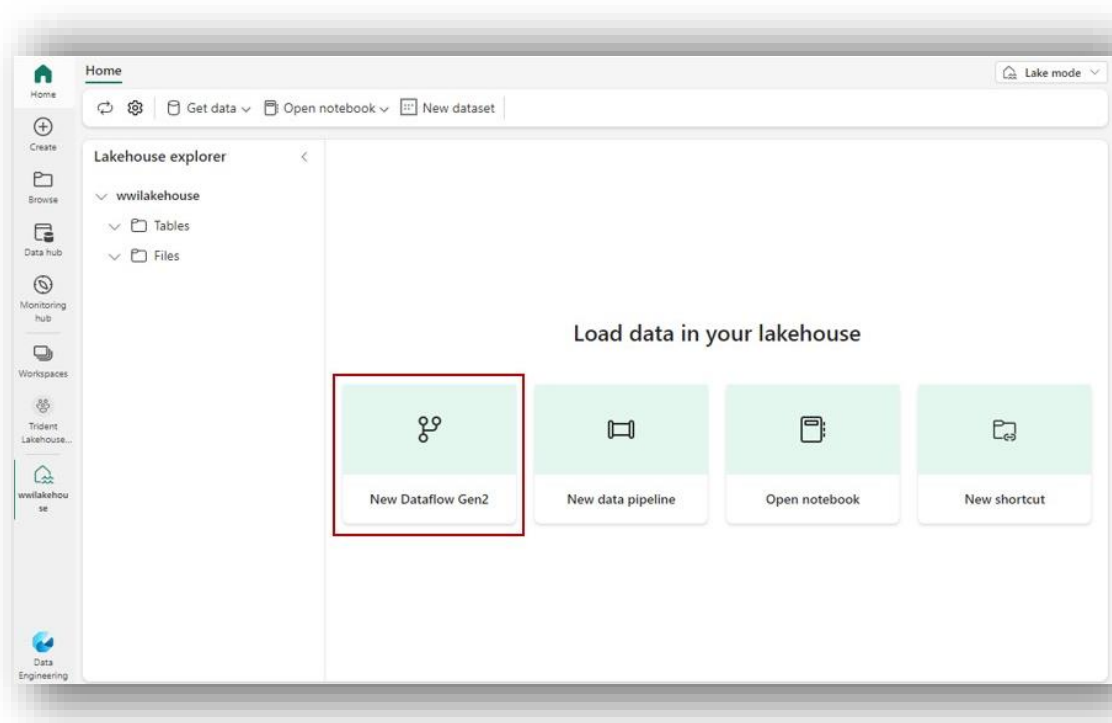
5. Enter **wwilakehouse** in the **Name** box.



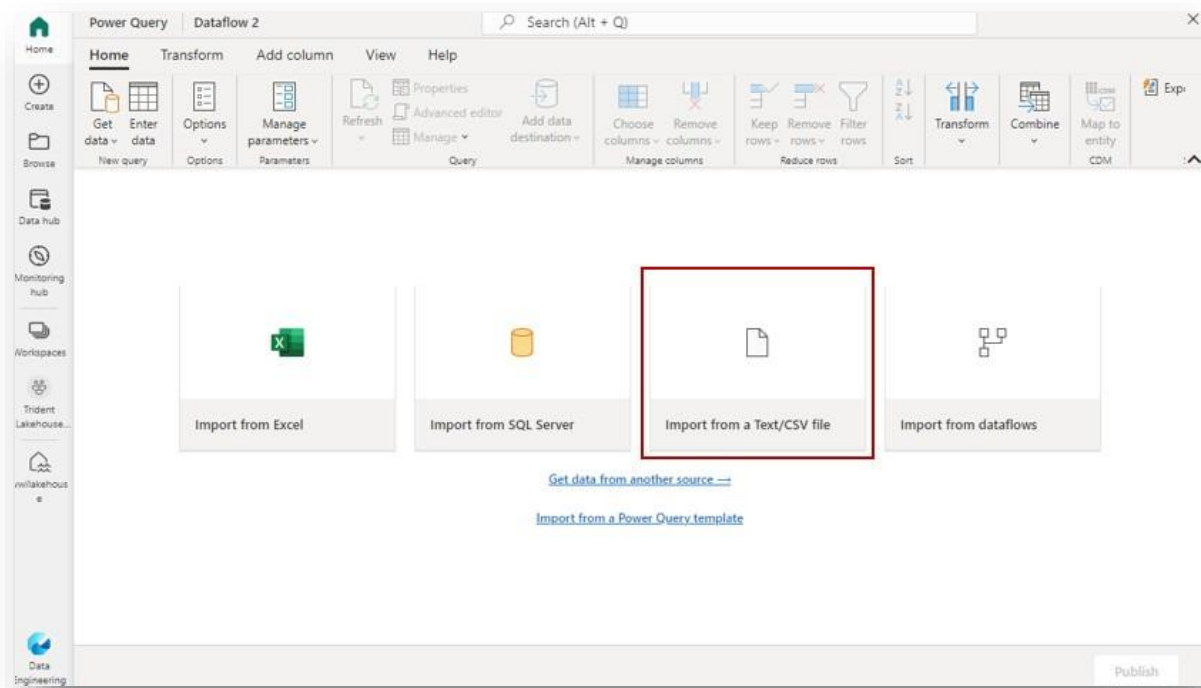
6. Click **Create**. The new lakehouse will be created and automatically opened.

Data Ingestion

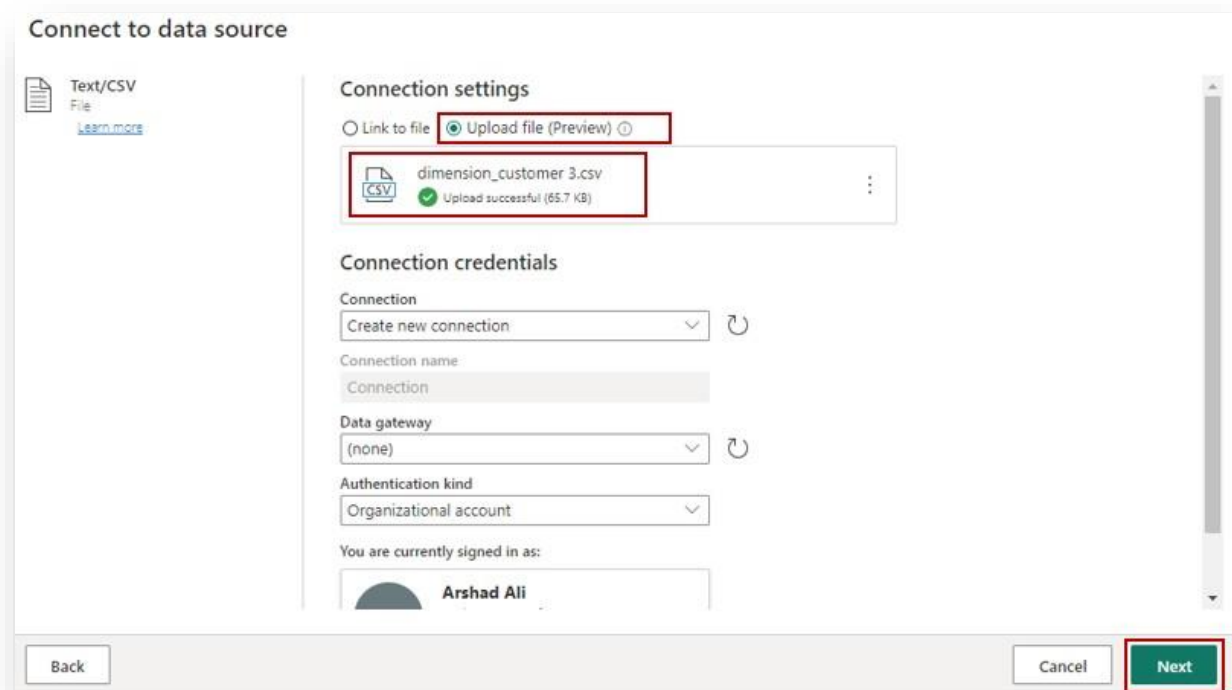
1. Download the 'dimension_customer.csv' file found in Data folder.
2. In the lakehouse view, you will see options to load data into lakehouse. Click on **New Dataflow Gen2**.



3. On the new dataflow page, click on **Import from a Text/CSV file**.



4. On the **Connect to data source** wizard, click on **Upload file** radio button and then drag and drop the data file that you downloaded in step 1 of this module. Once the file is uploaded, click **Next**.



5. Clicking **Next** on the previous screen will open the **Preview file data** page, click on **Create** to proceed and return back to dataflow canvas.

Preview file data

URL: https://microsoft-my.sharepoint.com/personal/arali_microsoft_com/Documents/Apps/Microsoft Power Query/Uploaded Files/dimension_customer 7.csv

File origin: 65001: Unicode (UTF-8) | Delimiter: Comma | Data type detection: Based on first 200 rows

CustomerKey	WWCustomerID	Customer	BillToCustomer	Category	BuyingGroup	PrimaryContact	PostalCode	ValidFrom	ValidTo
0	0	Unknown	N/A	N/A	N/A	N/A	N/A	1/1/2013, 12:00:00 AM	1/1/2025
234	433	Wingtip Toys (Baiko, OK)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Adirake Saenamuang	90584	1/1/2013, 12:00:00 AM	1/1/2025
235	434	Wingtip Toys (Lime Lake, NY)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Tarja Penttila	90667	1/1/2013, 12:00:00 AM	1/1/2025
236	435	Wingtip Toys (Teutopolis, IL)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Devendranath Huq	90269	1/1/2013, 12:00:00 AM	1/1/2025
237	436	Wingtip Toys (Gargatha, VA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Enes Olsson	90752	1/1/2013, 12:00:00 AM	1/1/2025
238	437	Wingtip Toys (Cadogan, PA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Vanja Grigic	90357	1/1/2013, 12:00:00 AM	1/1/2025
239	438	Wingtip Toys (Lucasville, OH)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Dayaram Mishra	90391	1/1/2013, 12:00:00 AM	1/1/2025
240	439	Wingtip Toys (Cylon, WI)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Mina Omidzadeh	90266	1/1/2013, 12:00:00 AM	1/1/2025
241	440	Wingtip Toys (Asher, OK)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Kadir Usenuly	90683	1/1/2013, 12:00:00 AM	1/1/2025
242	441	Wingtip Toys (Keosauqua, IA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Kajsa Jakobsson	90281	1/1/2013, 12:00:00 AM	1/1/2025
243	442	Wingtip Toys (Homer City, PA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Rahul Ghate	90713	1/1/2013, 12:00:00 AM	1/1/2025
244	443	Wingtip Toys (Benville, MI)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Lien Binh	90628	1/1/2013, 12:00:00 AM	1/1/2025
245	444	Wingtip Toys (Tea, SD)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Divyendu Chakraborty	90641	1/1/2013, 12:00:00 AM	1/1/2025
246	445	Wingtip Toys (Dacono, CO)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	An Dung Ngo	90289	1/1/2013, 12:00:00 AM	1/1/2025
247	446	Wingtip Toys (Saint Landry, LA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Chetana Kamath	90033	1/1/2013, 12:00:00 AM	1/1/2025
248	447	Wingtip Toys (Coin, IA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Premwadee Saowaluk	90458	1/1/2013, 12:00:00 AM	1/1/2025
249	448	Wingtip Toys (Salt Wells, NV)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Jae-Hwa Min	90058	1/1/2013, 12:00:00 AM	1/1/2025
250	449	Wingtip Toys (Delray, WV)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Catalina Nechita	90317	1/1/2013, 12:00:00 AM	1/1/2025
251	450	Wingtip Toys (Mount Summit, IN)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Constanza Laureano	90142	1/1/2013, 12:00:00 AM	1/1/2025

Back Cancel Add table using examples **Create**

6. In the **Name** property of the query settings pane, enter **dimension_customer**. Next, select the gear in the **Data destination** property at the bottom of the query settings pane.

Power Query

Search (Alt + Q)

Home Transform Add column View Help

Get Enter data Manage connections Options Refresh Advanced editor Manage Query Add data Choose Remove Keep Remove Filter Suggested Split Group Data type Whole number Merge queries Append queries Map to entity Copilot Export template

Queries [1] dimension_customer

Table.TransformColumnTypes(*Promoted headers*, {"CustomerKey", Int64.Type}, {"CustomerID", Int64.Type}, {"Customer", type text}, {"BillToCustomer", type text}, {"Category", type text})

CustomerKey	WWCustomerID	Customer	BillToCustomer	Category	BuyingGroup	PrimaryContact	PostalCode	ValidFrom	ValidTo
0	0	Unknown	N/A	N/A	N/A	N/A	N/A	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
234	433	Wingtip Toys (Baiko, OK)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Adirake Saenamuang	90584	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
235	434	Wingtip Toys (Lime Lake, NY)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Tarja Penttila	90667	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
236	435	Wingtip Toys (Teutopolis, IL)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Devendranath Huq	90269	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
237	436	Wingtip Toys (Gargatha, VA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Enes Olsson	90752	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
238	437	Wingtip Toys (Cadogan, PA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Vanja Grigic	90357	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
239	438	Wingtip Toys (Lucasville, OH)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Dayaram Mishra	90391	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
240	439	Wingtip Toys (Cylon, WI)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Mina Omidzadeh	90266	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
241	440	Wingtip Toys (Asher, OK)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Kadir Usenuly	90683	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
242	441	Wingtip Toys (Keosauqua, IA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Kajsa Jakobsson	90281	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
243	442	Wingtip Toys (Homer City, PA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Rahul Ghate	90713	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
244	443	Wingtip Toys (Benville, MI)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Lien Binh	90628	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
245	444	Wingtip Toys (Tea, SD)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Divyendu Chakraborty	90641	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
246	445	Wingtip Toys (Dacono, CO)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	An Dung Ngo	90289	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
247	446	Wingtip Toys (Saint Landry, LA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Chetana Kamath	90033	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
248	447	Wingtip Toys (Coin, IA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Premwadee Saowaluk	90458	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
249	448	Wingtip Toys (Salt Wells, NV)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Jae-Hwa Min	90058	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
250	449	Wingtip Toys (Delray, WV)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Catalina Nechita	90317	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
251	450	Wingtip Toys (Mount Summit, IN)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Constanza Laureano	90142	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
252	451	Wingtip Toys (Smart, WI)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Aakriti Bhandari	90071	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
253	452	Wingtip Toys (Lake Davis, CA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Himanshi Prabhakar	90719	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
254	453	Wingtip Toys (Staradovka, IL)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Bryan Harris	90739	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
255	454	Wingtip Toys (Concordia, PR)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Hani Kask	90486	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
256	455	Wingtip Toys (Tift Heights, CA)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Ike Sanderl	90304	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
257	456	Wingtip Toys (West Hempstead, NY)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Sushant Matondkar	90799	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
258	457	Wingtip Toys (Portales, NM)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Besavanna Anantham	90304	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
259	458	Wingtip Toys (Coker, AL)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Miguel Peaz	90364	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
260	459	Wingtip Toys (Haysville, OH)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Tobias Gruber	90519	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
261	460	Wingtip Toys (Mayville, NM)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Danica Dumora	90741	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM
262	461	Wingtip Toys (Oak, TX)	Wingtip Toys (Head Office)	Novelty Shop	Wingtip Toys	Ganesh Muthayee	90707	1/1/2013, 12:00:00 AM	1/1/2025, 12:00:00 AM

Completed (7.11 s) Column: 11 Rows: 99+

Query settings

Name: dimension_customer

Entity type: Custom

Applied steps: Promoted, Changed

Data destination: Lakehouse

Publish

7. If necessary, on the **Connect to data destination** screen, sign into your account. Click **Next**.
8. Navigate to the **wwilakehouse** in your workspace.
9. If the **dimension_customer** table does not exist, select the **New table** setting and enter the Table name of **dimension_customer**. If the table already exists, select the **Existing table** setting and select **dimension_customer** from the table list in the object explorer. Select **Next**.

Note – UI adds <space>Number at the end of the table name by default. Table names must be lower case and must not contain space. Please name it appropriately and remove any space from the table name.

The screenshot shows a 'Choose destination target' dialog box. At the top, there are two radio buttons: 'New table' (selected) and 'Existing table'. Below this is a search bar and a 'Display options' dropdown. On the left, a tree view shows a workspace with folders like 'Fabric Lakehouse Tuto...' and 'Dataflows/Staging/Lakeh...', and a selected item 'wwilakehouse'. On the right, a message states 'A new table will be created in wwilakehouse'. Below this, a 'Table name' field contains the text 'dimension_customer'. At the bottom, there are 'Back', 'Cancel', and 'Next' buttons. The 'Next' button is highlighted with a red box.

10. On the **Choose destination settings**, ensure **Use automatic settings** is enabled and click on **Save Settings**.

Choose destination settings

Use automatic settings

Column mapping

Source	Source type	Destination	Destination type
<input type="checkbox"/> CustomerKey	1 ² 3 Whole number	CustomerKey	Whole number
<input type="checkbox"/> WWICustomerID	1 ² 3 Whole number	WWICustomerID	Whole number
<input type="checkbox"/> Customer	A ^B C Text	Customer	Text
<input type="checkbox"/> BillToCustomer	A ^B C Text	BillToCustomer	Text
<input type="checkbox"/> Category	A ^B C Text	Category	Text
<input type="checkbox"/> BuyingGroup	A ^B C Text	BuyingGroup	Text
<input type="checkbox"/> PrimaryContact	A ^B C Text	PrimaryContact	Text
<input type="checkbox"/> PostalCode	A ^B C Text	PostalCode	Text
<input type="checkbox"/> ValidFrom	Date/Time	ValidFrom	Date/Time

Back

Cancel

Save settings

11. This will return you to the canvas of the dataflow. You can easily and quickly transform the data based on your business requirements using this nice and intuitive graphical user interface. For the purpose of this module, we will not make any changes here. To proceed, click on **Publish** at bottom right of the screen.

Query settings >

▼ Properties

Name

dimension_customer

Entity type ⓘ
Custom

▼ Applied steps

Source

Promoted h...

ABC 123 Changed co...

▼ Data destination

Lakehouse

Step


Publish

14. Once the dataflow's refresh is completed, you can go to the lakehouse, and you will notice **dimension_customer** delta table has been created. When you click on it, you should be able to preview its data. Further, you can use the SQL endpoint of the lakehouse to query the data with SQL statements in warehouse mode. Click on **SQL endpoint** under **Lake mode** on top right of the screen.

The screenshot displays the Microsoft Fabric Lakehouse interface. On the left, the 'Explorer' pane shows the 'wwilakehouse' containing a 'Tables' folder with the 'dimension_customer' table highlighted. The main area shows a preview of the 'dimension_customer' table with 403 rows. A red box highlights the 'SQL analytics endpoint' in the top right corner, which is used for querying data using SQL.

CustomerKey	WWCustomerID	Customer	BillToCustomer	Category	BuyingGroup	PrimaryContact	PostalCode	ValidFrom	ValidTo	LineageKey
0	0	Unknown	N/A	N/A	N/A	N/A	N/A	1/1/2013 12:00:00...	1/1/2025 12:00:00...	0
102	102	Tailspin Toys (Fieldbr...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Tea Koppel	90205	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
103	103	Tailspin Toys (Kilvest...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Naseem Radan	90130	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
104	104	Tailspin Toys (Wallagr...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Laboni Deb	90579	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
105	105	Tailspin Toys (Tomnol...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Sung-Hwan Hwa...	90400	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
106	106	Tailspin Toys (Tumaca...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Shiva Pipalia	90662	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
107	107	Tailspin Toys (Glen Av...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Karie Mercier	90782	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
108	108	Tailspin Toys (Bernie, ...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Bhanu Thota	90045	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
109	109	Tailspin Toys (South L...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Ae-Cha Joo	90247	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
110	110	Tailspin Toys (North C...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Dinara Saparkyzy	90792	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2
111	111	Tailspin Toys (Oniole ...	Tailspin Toys (He...	Novelty Shop	Kids Toys	Adam Dvorak	90436	1/1/2013 12:00:00...	1/1/2025 12:00:00...	2


15. In the warehouse mode, you can click on **dimension_customer** table to preview its data and/or click on **New SQL query** to write your SQL statements.



16. Here is a sample query to aggregate the row count based on buyinggroup column of the **dimension_customer** table and its output. SQL query files are saved automatically for future references, and you can rename or delete these files appropriately based on your need.

To run the script, click on the **Run** icon at the top of the script file.

```
SELECT BuyingGroup, Count(*) AS Total
FROM dimension_customer
GROUP BY BuyingGroup
```

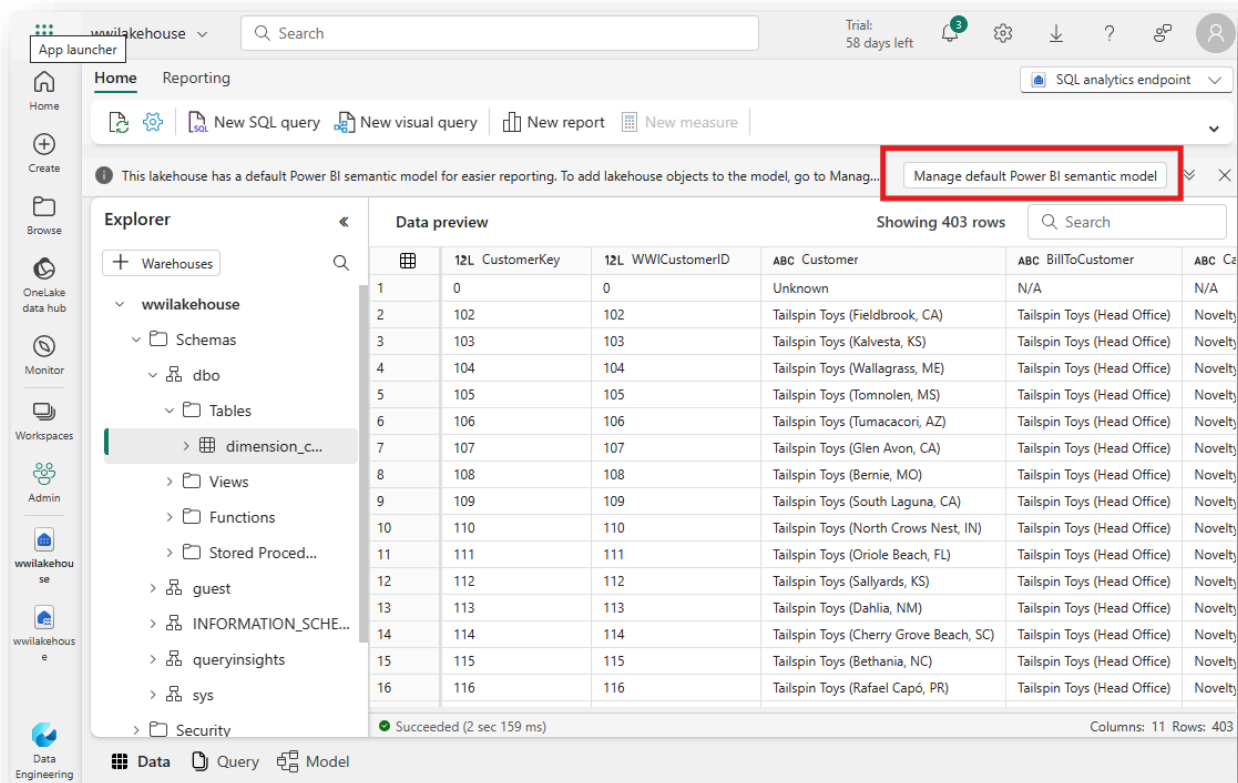


Add table to the semantic model

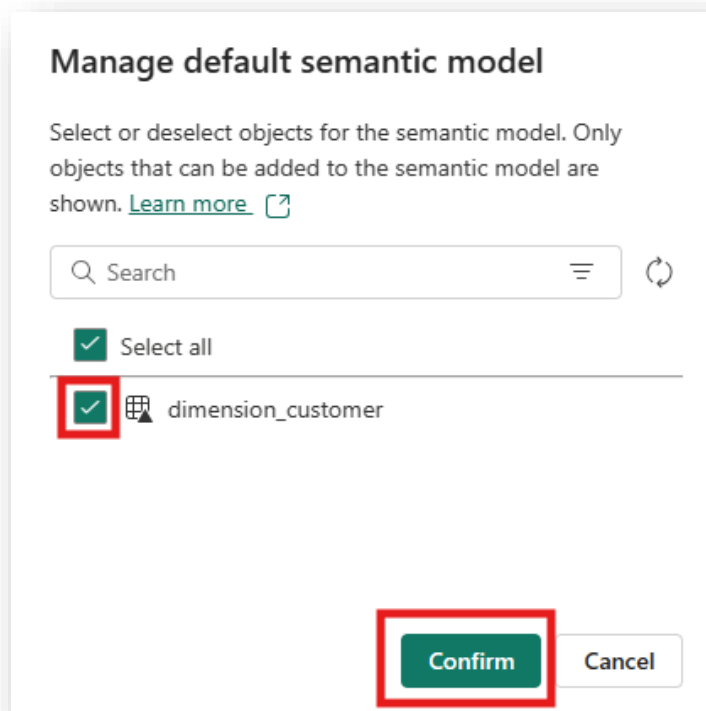
In Microsoft Fabric, Power BI semantic models are a logical description of an analytical domain, with metrics, business friendly terminology, and representation, to enable deeper analysis. This semantic model is typically a star schema with facts that represent a domain, and dimensions that allow you to analyze, or slice and dice the domain to drill down, filter, and calculate different analyses.

Now that you've ingested data into a lakehouse table, we'll add that table to the **default semantic model** for the lakehouse.

1. From the SQL analytics Data view of the dimension_customer table, click **Manage default Power BI semantic model** button on the upper right.

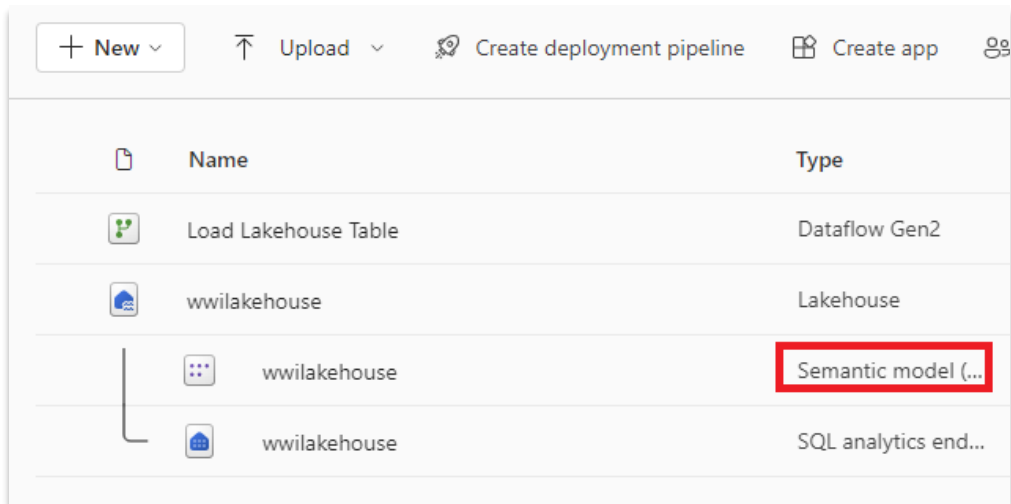






- The **Manage default semantic model** dialog will appear. Select the `dimension_customer` table and click **Confirm**.



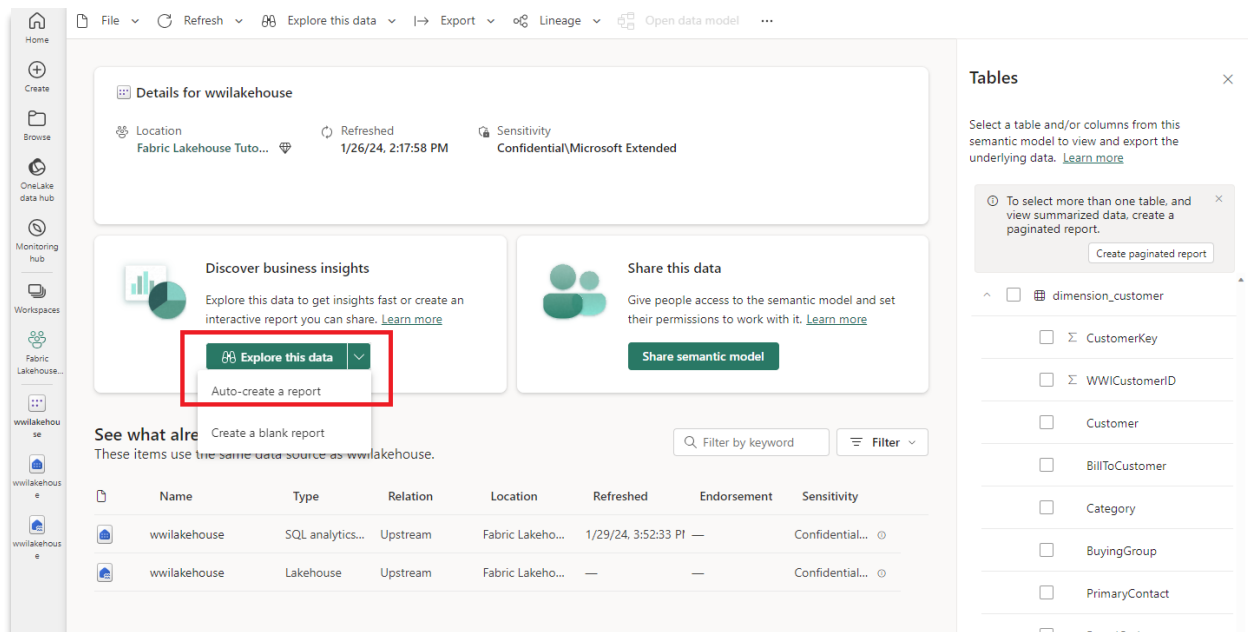
Building a report

1. In the artifact view of the workspace, click on **wwilakehouse** default semantic model, which gets created automatically with the same name of the lakehouse when you create a lakehouse.



+ New ▾ ↶ Upload ▾ 🚀 Create deployment pipeline 📦 Create app 👤		
	Name	Type
	Load Lakehouse Table	Dataflow Gen2
	wwilakehouse	Lakehouse
	wwilakehouse	Semantic model (...)
	wwilakehouse	SQL analytics end...

- On the semantic model screen, you can view all the tables. You will have options to create reports either from scratch, paginated report or let Power BI do magic for you by automatically creating a report based on your data. For the purpose of this module, click on **Auto-create a report** under **Explore this data**. In the next module, we will create a report from scratch.



- Since the table is a dimension and there are no measures in it, Power BI smartly creates a measure for the row sum and aggregates it across different columns and creates different charts as below.

You can save this report for the future by clicking on **Save** button at the top ribbon and giving it a name. You can further make changes to this report to meet your requirement by including or excluding additional tables or columns.



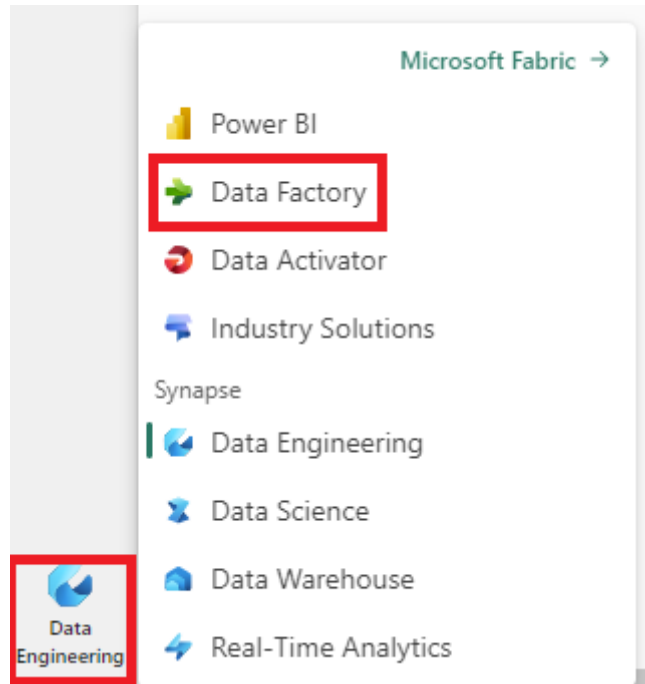
Module 2: Ingest, Prep and Analyze

This module is going to build on the work you completed in the previous module and ingest additional tables (dimensions and fact) of the dimensional model of Wide World Importers (WWI) as mentioned in the Introduction section of this document. Next, you will use notebooks with Spark runtime to transform and prepare the data. Finally, you will create Power BI data model and create a report from scratch.

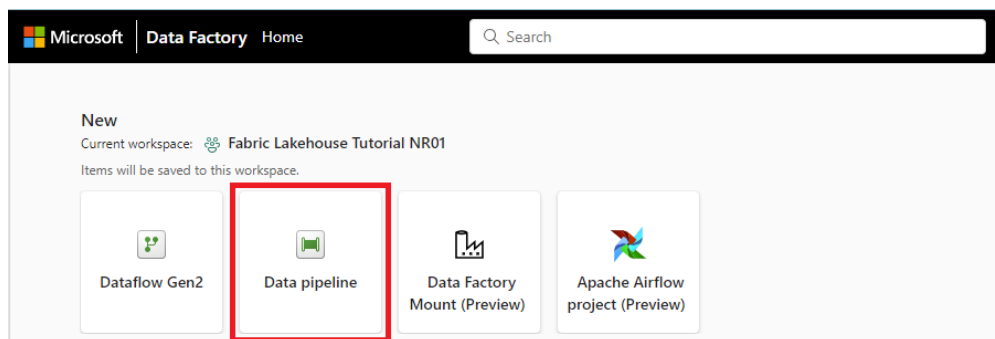
Data Ingestion

In this section, you will use **Copy data activity** of **Data Factory pipeline** to ingest sample data from source (Azure storage account) to the **Files** section of the lakehouse you created earlier.

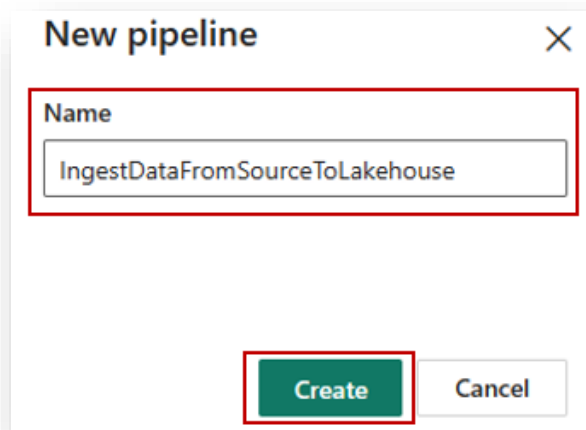
1. On the bottom left of the screen, click on the workload switcher, and then select **Data Factory**.



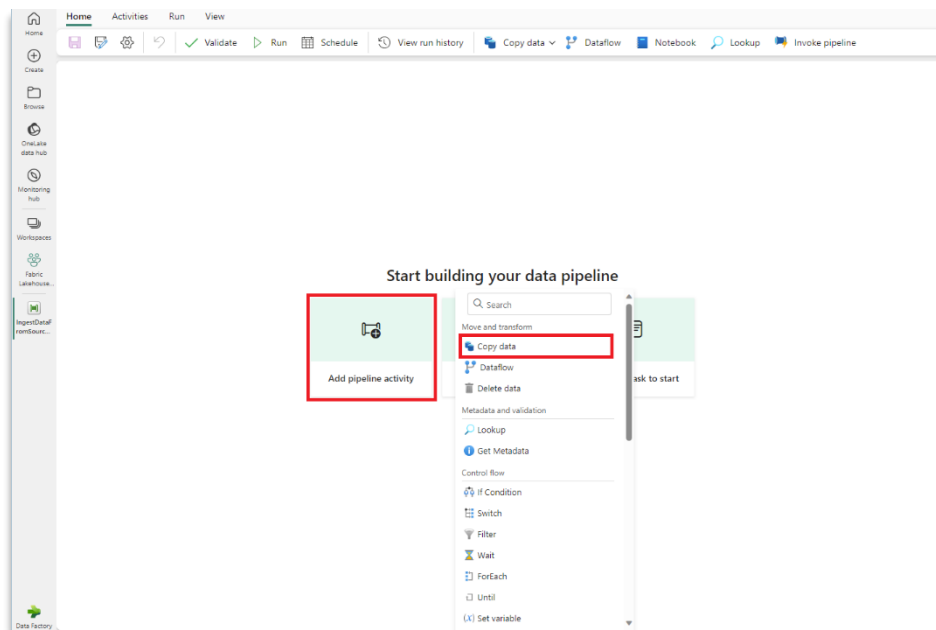
2. Click on **Data pipeline** under **New** to create a new data pipeline.



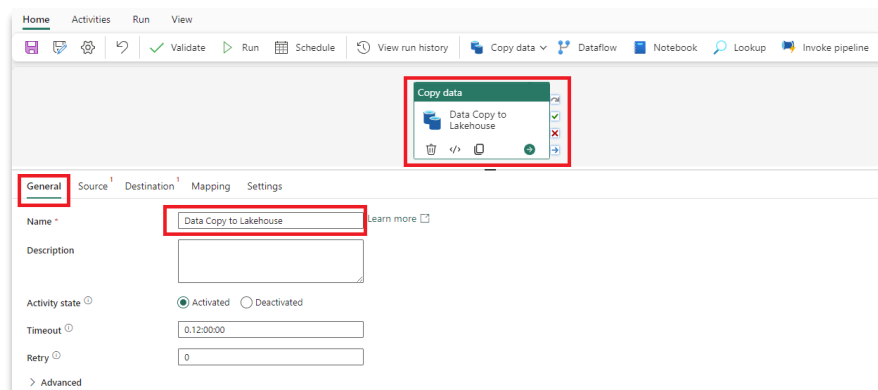
3. For the **New pipeline**, specify the name as **IngestDataFromSourceToLakehouse** and click on **Create**. This will create a new data factory pipeline and open it on the screen to work on it.




- On the newly created data factory pipeline, click **Add pipeline activity** to add an activity to the pipeline and click on **Copy data**. This will add copy data activity to the pipeline canvas.



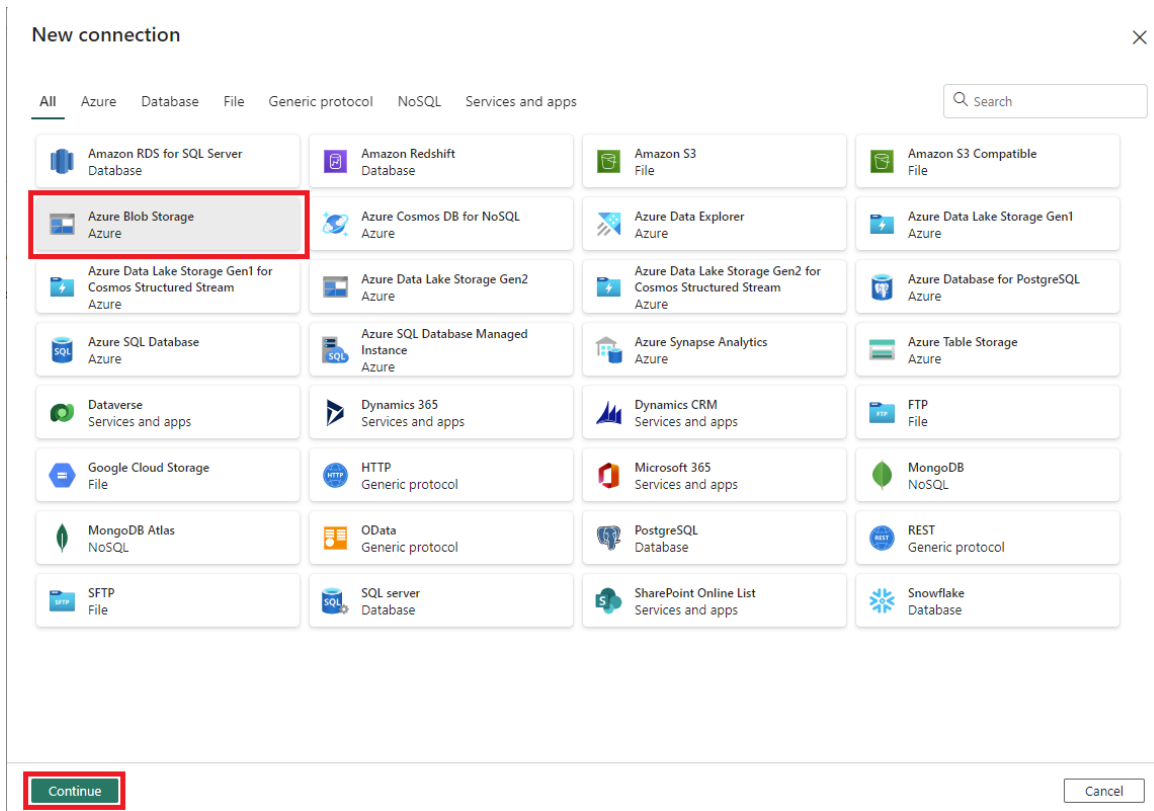
- Select the newly added copy data activity to the canvas, which will show activity properties at the bottom. Under **General** tab, specify the name for the copy data activity **Data Copy to Lakehouse**.



6. Under **Source** tab of the selected copy data activity, select **External** as **Data store type** and then click on **+ New** to create a new connection to data source.



7. For this tutorial, all the sample data is available in a public container of Azure blob storage so you will connect to this container to copy data from. On the **New connection wizard**, select **Azure Blob Storage** and then click on **Continue**.



8. On the next screen of the **New connection wizard**, please specify these details and click on **Create** to create the connection to the data source.

Account name or URI	https://azuresynapsestorage.blob.core.windows.net/sampledata
Connection	Create new connection
Connection name	wwisampledta
Authentication kind	Anonymous

New connection

Azure Blob Storage
Learn more

Connection settings

Account name or URL *
https://azuresynapsestorage.blob.core.windows.net/sampl...

Connection credentials

Connection
Create new connection

Connection name
wwisampledta

Authentication kind
Anonymous

Create Back Cancel

9. Once the new connection is created, you will return to **Source** tab of the copy data activity and newly created connection will be selected by default. Please specify these properties before moving to the destination settings.

Data store type	External
Connection	wwisampledta
File path type	File path
File path	Container name (first text box) – sampledata Directory name (second text box) – WideWorldImportersDW/parquet
Recursively	Checked
File Format	Binary

Home Activities Run View

Validate Run Schedule View run history Copy data Dataflow Notebook

Copy data

Data Copy to Lakehouse

General **Source** Destination Mapping Settings

Data store type
☐ Workspace ☒ External ☐ Sample dataset

Connection *
wwisampledta

File path type
☒ File path ☐ Prefix ☐ Wildcard file path ☐ List of files

File path *
sampledata / WideWorldImporters...


Recursively ☒

File format *
Binary

> Advanced

10. Under **Destination** tab of the selected copy data activity, please specify these properties:

Data store type	Workspace
Workspace data store type	Lakehouse
Lakehouse	Wwilakehouse
Root Folder	Files
File path	Directory name (first text box) – wwi-raw-data
File Format	Binary



11. At this time, you have completed configuring copy data activity. Click on **Save** button under **Home** to save all these details and click on **Run** to kick off execution of this pipeline and its activity. You can also schedule pipelines to refresh data at defined intervals to meet your business requirements, however for this module, we will run the pipeline once by clicking on **Run** button. This will trigger data copy from the underlying data source to the specified lakehouse and might take up to a minute to complete. You can monitor the execution of the pipeline and its activity under **Output** tab which appears when you click anywhere on the canvas. Optionally, you can click on the activity name **“Data Copy to Lakehouse”** to look at the details of data transfer.

The screenshot displays the Azure Data Factory user interface. At the top, the 'Run' button is highlighted with a red box. Below the canvas, the 'Output' tab is selected and highlighted with a red box. The 'Output' tab shows the pipeline run ID, status, and a table of activities. The table has columns for Activity name, Activity status, Run start, Duration, Input, and Output. The first row shows the activity 'Data Copy to Lakehouse' with a status of 'Succeeded'.

Activity name	Activity status	Run start	Duration	Input	Output
Data Copy to Lakehouse	Succeeded	1/29/2024, 6:02:35 PM	30s		

12. Once the data copy is completed, go to items view of the workspace and click on **wwilakehouse** to launch **Lakehouse explorer** of this selected lakehouse.

Fabric Lakehouse Tutorial NR01
This workspace contains all the items for the lakehouse tutorial.

Filter by keyword Filter

Name	Type	Owner	Refreshed	Next refresh	Endorsement	Sensitivity
IngestDataFromSourceToLakehouse	Data pipeline	Natalia Rodriguez ...	—	—	—	Confidential
Load Lakehouse Table	Dataflow Gen2	Natalia Rodriguez ...	1/29/24, 3:49:05 PM	N/A	—	Confidential
report1	Report	Fabric Lakehouse ...	1/26/24, 2:17:58 PM	—	—	Confidential
wwilakehouse	Lakehouse	Natalia Rodriguez ...	—	—	—	Confidential
wwilakehouse	Semantic model (...)	Fabric Lakehouse ...	1/26/24, 2:17:58 PM	N/A	—	Confidential
wwilakehouse	SQL analytics endpoint	Fabric Lakehouse ...	1/29/24, 3:52:33 PM	N/A	—	Confidential

13. In the **Lakehouse explorer** view, you would notice a new folder **wwi-raw-data** has been created and data for all the tables have been copied here.

Home Lakehouse

Get data New semantic model Open notebook

A SQL analytics endpoint for SQL querying and a default Power BI semantic model for faster reporting were created and will be updated with any tables added to the lakehouse.

Explorer

- wwilakehouse
 - Tables
 - Files
 - wwi-raw-data**
 - full
 - dimension_city
 - dimension_customer
 - dimension_date
 - dimension_employ...
 - dimension_stock_it...
 - fact_sale
 - fact_sale_1y_f...**
 - incremental
 - tables


Files > wwi-raw-data > full > **fact_sale_1y_full**

Name	Date modified	Type	Size
_SUCCESS	1/29/2024 6:02:5...	-	0 B
part-00000-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	22 MB
part-00001-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	26 MB
part-00002-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	19 MB
part-00003-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	19 MB
part-00004-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	33 MB
part-00005-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	20 MB
part-00006-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	36 MB
part-00007-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	23 MB
part-00008-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	24 MB
part-00009-ced648ca-e8c8-46e5-8526-5ca85d56e67e-c000....	1/29/2024 6:03:0...	PARQUET	24 MB

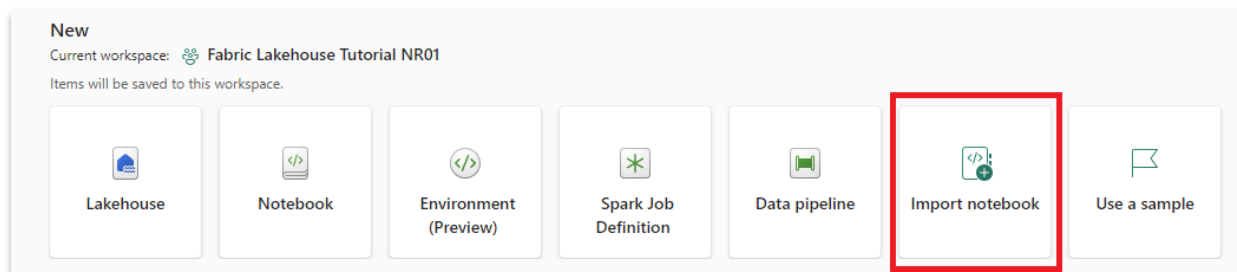
Data Preparation

Now since we have raw data already ingested from source to the **Files** section of the lakehouse, you can take this data, transform, and prepare it to create delta tables.


1. Download the set of notebooks found in the Scripts folder
2. From the workload switcher located at the bottom left of the screen, select **Data engineering**.



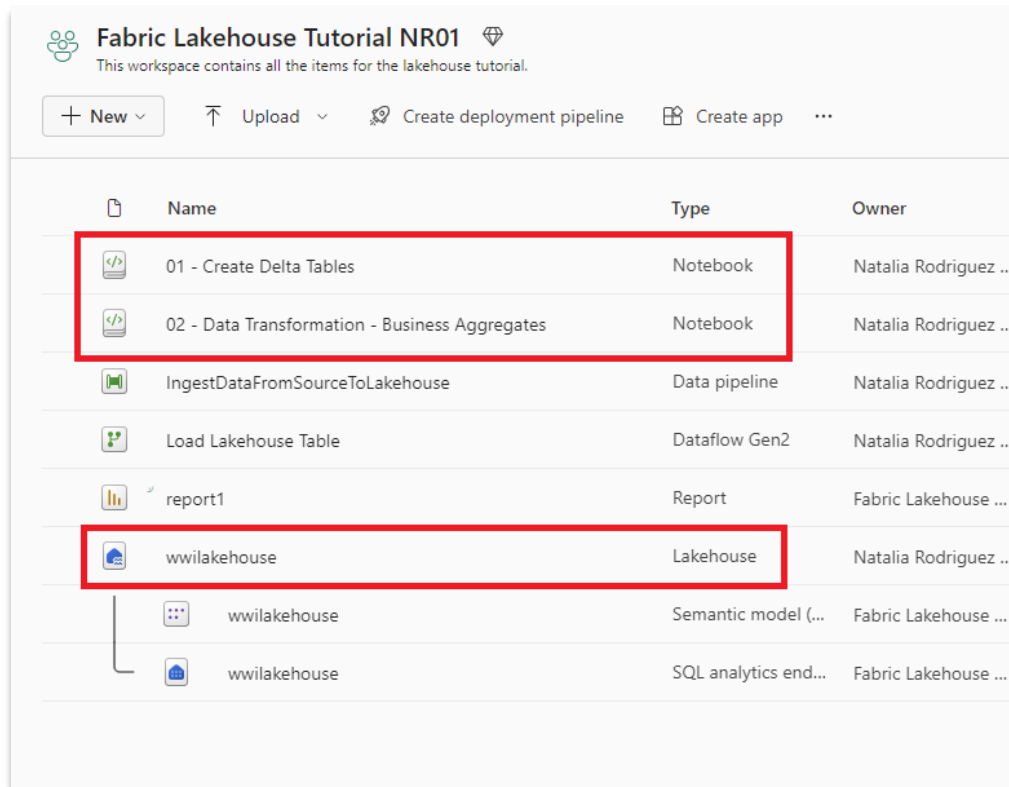
3. Select **Import notebook** from the **New** section at the top of the landing page.




4. Select **Upload** from the **Import status** pane that opens on the right side of the screen.
5. Select all the notebooks that were downloaded and/or unzipped in step 1 of this section.
6. Select **Open**. A notification indicating the status of the import will appear in the top right corner of the browser window.



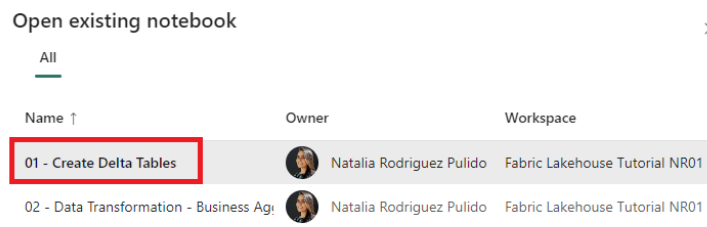
7. After the import of notebooks is successful, you can go to items view of the workspace and see these newly imported notebooks. Click on **wwilakehouse** lakehouse to open it.



8. Once the **wwilakehouse** lakehouse is opened, click on **Open notebook > Existing notebook** on the ribbon at the top.



9. From the list of existing notebooks, select the **01 - Create Delta Tables** notebook and click on **Open** button.



10. Once the notebook is opened, in the **Lakehouse explorer** you will notice the notebook is already linked to your opened lakehouse.

Note

Fabric provides these unique capabilities for writing optimized delta lake files:

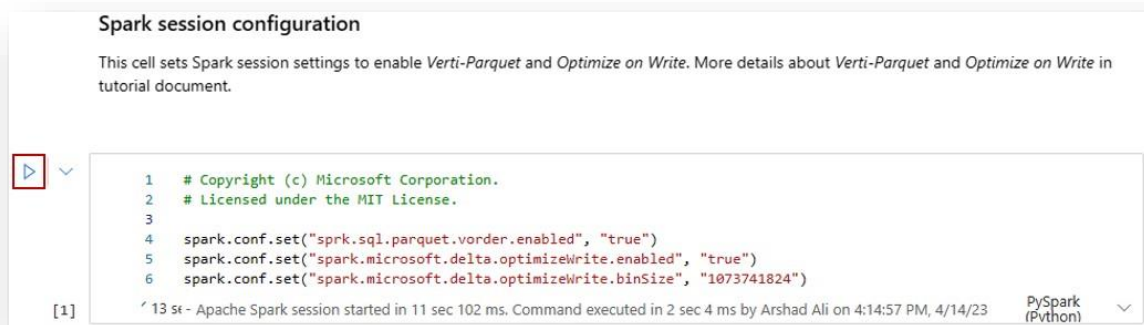
- V-Order – Fabric includes Microsoft’s unique V-Order IP. V-Order transparently optimizes the Delta Lake files in a way that is highly optimized by Fabric

compute engines, often resulting in 3x-4x compression improvement and up to 10x performance acceleration over Delta Lake files not optimized using VertiParquet while still maintaining full Delta Lake format compliance.

- [Optimize write](#) – Apache Spark performs most efficiently when using standardized larger file sizes. The relation between the file size, the number of files, the number of Spark workers and Spark's configurations play a critical role in performance. Ingestion workloads into Delta Lake tables may have the inherited characteristic of constantly writing lots of small files; this scenario is commonly known as the "small files problem". To overcome this problem, Spark in Fabric includes an Optimize Write feature that reduces the number of files written and aims to increase individual file size of the written data. It dynamically optimizes partitions while generating files with a default 128 MB size. The target file size may be changed per workload requirements using configurations.

11. Before you write data as delta lake tables in the **Tables** section of the lakehouse, you are going to use two features (**V-Order** and **Optimize Write**) of Fabric for optimized data writing and for subsequent better reading performance. To enable these features in your session you can set these configurations in the first cell of your notebook - eventually these features will be by default enabled for the Spark session.

To start execution, click **Run All** under **Home** at the top to start execution of the notebook and all its cells in the sequence or click **Run** icon on the left of the cell or hit **SHIFT+ENTER** while control is in the cell to execute code from that specific cell.



```
Spark session configuration

This cell sets Spark session settings to enable Verti-Parquet and Optimize on Write. More details about Verti-Parquet and Optimize on Write in tutorial document.

1 # Copyright (c) Microsoft Corporation.
2 # Licensed under the MIT License.
3
4 spark.conf.set('spark.sql.parquet.vorder.enabled', 'true')
5 spark.conf.set('spark.microsoft.delta.optimizeWrite.enabled', 'true')
6 spark.conf.set('spark.microsoft.delta.optimizeWrite.binSize', '1073741824')

[1] 13 s - Apache Spark session started in 11 sec 102 ms. Command executed in 2 sec 4 ms by Arshad Ali on 4:14:57 PM, 4/14/23 PySpark (Python)
```

The first thing you would have noticed when executing this cell is that you didn't have to specify the underlying Spark pool or cluster details as Fabric provides it through the concepts of Live Pool i.e., every Fabric workspace comes pre-wired with a default Spark pool, called Live Pool. This means that when users create notebooks, they don't have to worry about specifying any Spark configurations or cluster details (or something like that) and when they execute their first command in the notebooks the live pool kicks in and is up in a few seconds after establishing their Spark session and starts executing the code in the cell. Subsequent code execution is almost instantaneous in this notebook as long as the Spark session is active.

12. Next, you will read raw data from the **Files** section of the lakehouse, add additional columns for different date parts as part of the transformation. Finally, you will use partitionBy Spark API to partition the data before writing it as delta table based on the newly created data part columns (Year and Quarter).

Fact - Sale

This cell reads raw data from the *Files* section of the lakehouse, adds additional columns for different date parts and the same information is being used to create partitioned fact delta table.

```
1 from pyspark.sql.functions import col, year, month, quarter
2
3 table_name = 'fact_sale'
4
5 df = spark.read.format("parquet").load('Files/wwi-raw-data/full/fact_sale_1y_full')
6 df = df.withColumn('Year', year(col("InvoiceDateKey")))
7 df = df.withColumn('Quarter', quarter(col("InvoiceDateKey")))
8 df = df.withColumn('Month', month(col("InvoiceDateKey")))
9
10 df.write.mode("overwrite").format("delta").partitionBy("Year", "Quarter").save("Tables/" + table_name)
```

✓ - Command executed in 55 sec 644 ms by Arshad Ali on 11:03:21 AM, 4/14/23

PySpark (Python) ▾

```
from pyspark.sql.functions import col, year, month, quarter
```

```
table_name = 'fact_sale'
```

```
df = spark.read.format("parquet").load('Files/wwi-raw-data/full/fact_sale_1y_full')
df = df.withColumn('Year', year(col("InvoiceDateKey")))
df = df.withColumn('Quarter', quarter(col("InvoiceDateKey")))
df = df.withColumn('Month', month(col("InvoiceDateKey")))
```

```
df.write.mode("overwrite").format("delta").partitionBy("Year", "Quarter").save("Tables/" + table_name)
```

13. After fact data load, you can move on to loading data for the rest of the dimensions. The following cell creates a function to read raw data from the **Files** section of the lakehouse for each of table names passed as a parameter. Next, it creates a list of dimension tables. Finally, it has a for loop to loop through the list of tables and call created function with each table name as parameter to read data for that specific table and create delta table respectively.

Dimensions

This cell creates a function to read raw data from the *Files* section of the lakehouse for the table name passed as a parameter. Next, it creates a list of dimension tables. Finally, it has a *for loop* to loop through the list of tables and call above function with each table name as parameter to read data for that specific table and create delta table.

```
1 from pyspark.sql.types import *
2
3 def loadFullDataFromSource(table_name):
4     df = spark.read.format("parquet").load('Files/wwi-raw-data/full/' + table_name)
5     df = df.select([c for c in df.columns if c != 'Photo'])
6     df.write.mode("overwrite").format("delta").save("Tables/" + table_name)
7
8 full_tables = [
9     'dimension_city',
10    'dimension_date',
11    'dimension_employee',
12    'dimension_stock_item'
13 ]
14
15 for table in full_tables:
16     loadFullDataFromSource(table)
```

```

from pyspark.sql.types import *


def loadFullDataFromSource(table_name):
    df = spark.read.format("parquet").load('Files/wwi-raw-data/full/' +
        table_name)
    df = df.select([c for c in df.columns if c != 'Photo'])
    df.write.mode("overwrite").format("delta").save("Tables/" + table_name)

full_tables = [
    'dimension_city',
    'dimension_date',
    'dimension_employee',
    'dimension_stock_item'
]

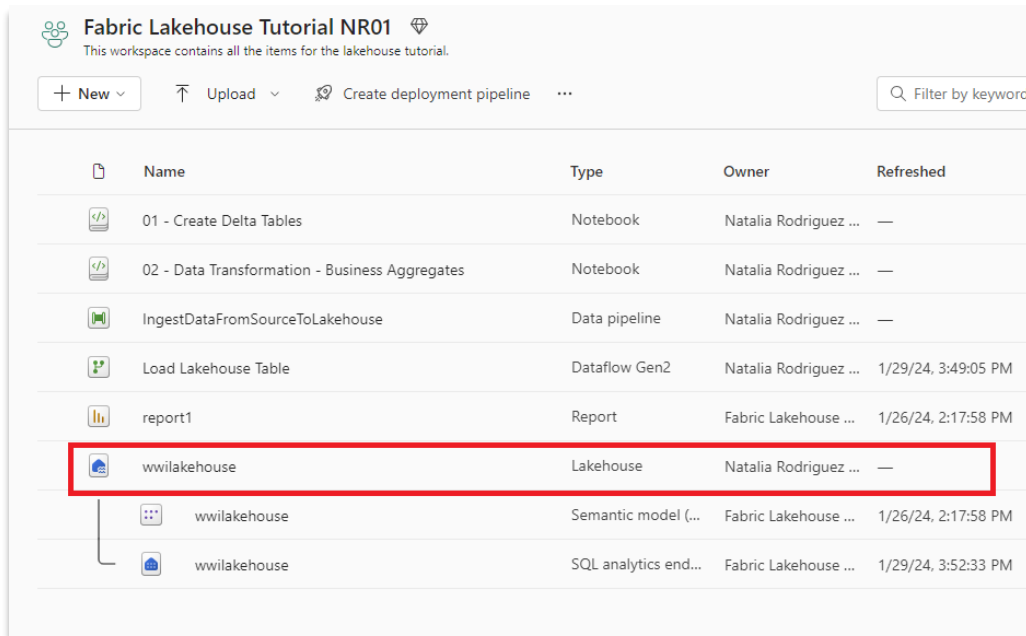
for table in full_tables:
    loadFullDataFromSource(table)

```

14. You can validate created tables by right clicking and selecting refresh on **wwilakehouse** lakehouse and you will notice these created tables.



15. Please go the items view of the workspace again and click on **wwilakehouse** lakehouse to open it.



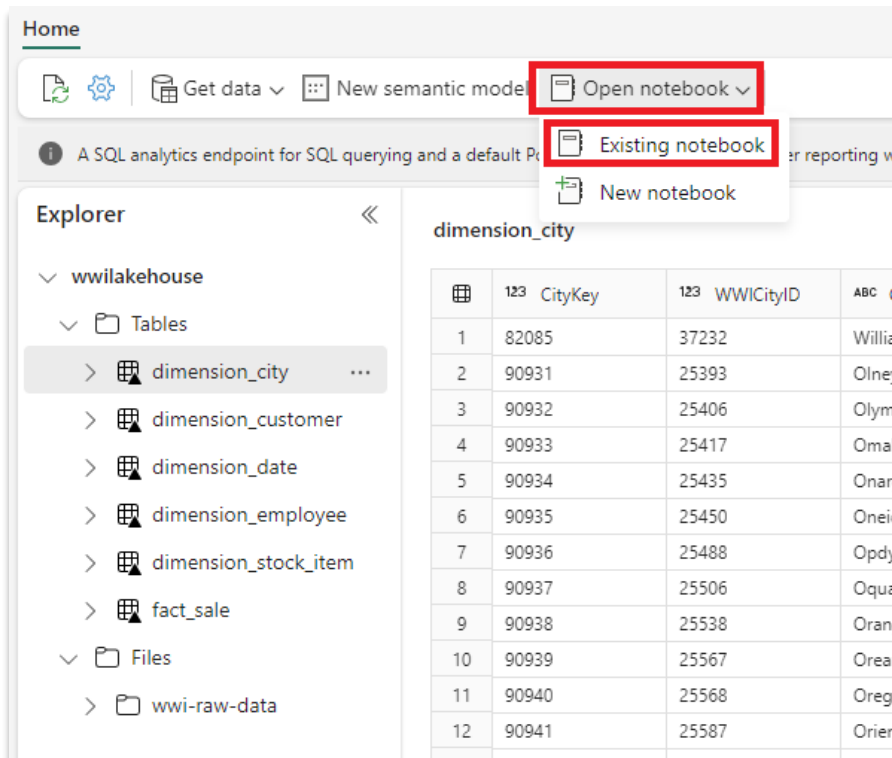
Fabric Lakehouse Tutorial NR01

This workspace contains all the items for the lakehouse tutorial.

+ New Upload Create deployment pipeline Filter by keyword

Name	Type	Owner	Refreshed
01 - Create Delta Tables	Notebook	Natalia Rodriguez ...	—
02 - Data Transformation - Business Aggregates	Notebook	Natalia Rodriguez ...	—
IngestDataFromSourceToLakehouse	Data pipeline	Natalia Rodriguez ...	—
Load Lakehouse Table	Dataflow Gen2	Natalia Rodriguez ...	1/29/24, 3:49:05 PM
report1	Report	Fabric Lakehouse ...	1/26/24, 2:17:58 PM
wwilakehouse	Lakehouse	Natalia Rodriguez ...	—
wwilakehouse	Semantic model (...)	Fabric Lakehouse ...	1/26/24, 2:17:58 PM
wwilakehouse	SQL analytics end...	Fabric Lakehouse ...	1/29/24, 3:52:33 PM

16. Now, open the second notebook. In the lakehouse view, click on **Open notebook > Existing notebook** on the ribbon at the top.



Home

Get data New semantic model Open notebook Existing notebook New notebook

A SQL analytics endpoint for SQL querying and a default P... er reporting w

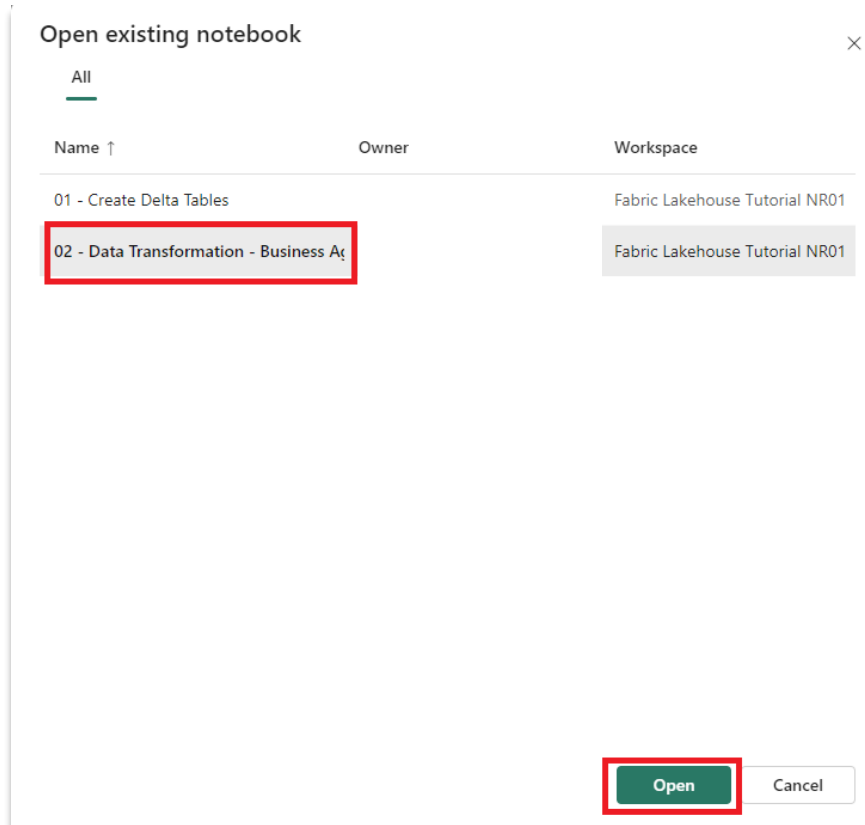
Explorer

- wwilakehouse
 - Tables
 - dimension_city
 - dimension_customer
 - dimension_date
 - dimension_employee
 - dimension_stock_item
 - fact_sale
 - Files
 - wwi-raw-data

dimension_city

	123	CityKey	123	WWCityID	ABC
1	82085		37232		Willia
2	90931		25393		Olney
3	90932		25406		Olym
4	90933		25417		Oma
5	90934		25435		Onar
6	90935		25450		Onei
7	90936		25488		Opdy
8	90937		25506		Oqua
9	90938		25538		Oran
10	90939		25567		Orea
11	90940		25568		Oreg
12	90941		25587		Orien

17. From the list of existing notebooks, select the **02 - Data Transformation - Business** notebook to open it.
18. Once the notebook is opened, in the **Lakehouse explorer** you will notice the notebook is already linked to your opened lakehouse.
19. An organization might have a group of data engineers working with Scala/Python while there might be other groups of data engineers preferring to work with SQL (Spark SQL or T-SQL) and all these while working on the



same copy of the data. Fabric makes it possible for these different groups, with varied experience and preference, to work and collaborate.

You are going to learn two different approaches, as below, to transform and generate business aggregates, and as a developer you can pick the one suitable for you or mix and match these approaches based on your preference without compromising on the performance:

- Approach #1 – Use PySpark to join and aggregates data for generating business aggregates. This approach would be preferable to someone with a programming (Python or PySpark) background.
- Approach #2 – Use Spark SQL to join and aggregates data for generating business aggregates. This approach would be preferable to someone with SQL background, transitioning to Spark.

20. **Approach #1** – Use PySpark to join and aggregates data for generating business aggregates. In the below code, you are creating three different Spark dataframes, each referencing an existing delta table. Then, you are joining these tables using the dataframes, doing group by to generate aggregation, renaming few of the columns and finally writing it as delta table in the **Tables** section of the lakehouse to persist with the data.

Approach #1 - sale_by_date_city

In this cell, you are creating three different Spark dataframes, each referencing an existing delta table.

```
1 df_fact_sale = spark.read.table("wwilakehouse.fact_sale")
2 df_dimension_date = spark.read.table("wwilakehouse.dimension_date")
3 df_dimension_city = spark.read.table("wwilakehouse.dimension_city")
```

✓ 59 sec - Command executed in 59 sec 175 ms by Arshad Ali on 4:15:57 PM, 4/14/23

PySpark (Python) ▾

> Spark jobs (9 of 9 succeeded) ...

+ Code + Markdown

In this cell, you are joining these tables using the dataframes created earlier, doing group by to generate aggregation, renaming few of the columns and finally writing it as delta table in the Tables section of the lakehouse.

+ Code + Markdown

```
1 sale_by_date_city = df_fact_sale.alias("sale") \
2 .join(df_dimension_date.alias("date"), df_fact_sale.InvoiceDateKey == df_dimension_date.Date, "inner") \
3 .join(df_dimension_city.alias("city"), df_fact_sale.CityKey == df_dimension_city.CityKey, "inner") \
4 .select("date.Date", "date.CalendarMonthLabel", "date.Day", "date.ShortMonth", "date.CalendarYear", "city.City", "city.StateProvince",
5 "city.SalesTerritory", "sale.TotalExcludingTax", "sale.TaxAmount", "sale.TotalIncludingTax", "sale.Profit") \
6 .groupBy("date.Date", "date.CalendarMonthLabel", "date.Day", "date.ShortMonth", "date.CalendarYear", "city.City", "city.StateProvince",
7 "city.SalesTerritory") \
8 .sum("sale.TotalExcludingTax", "sale.TaxAmount", "sale.TotalIncludingTax", "sale.Profit") \
9 .withColumnRenamed("sum(TotalExcludingTax)", "SumOfTotalExcludingTax") \
10 .withColumnRenamed("sum(TaxAmount)", "SumOfTaxAmount") \
11 .withColumnRenamed("sum(TotalIncludingTax)", "SumOfTotalIncludingTax") \
12 .withColumnRenamed("sum(Profit)", "SumOfProfit") \
13 .orderBy("date.Date", "city.StateProvince", "city.City")
```

✓ 19 sec - Command executed in 19 sec 85 ms by Arshad Ali on 4:16:16 PM, 4/14/23

PySpark (Python) ▾

> Spark jobs (12 of 12 succeeded) ...

```
df_fact_sale = spark.read.table("wwilakehouse.fact_sale")
df_dimension_date = spark.read.table("wwilakehouse.dimension_date")
df_dimension_city = spark.read.table("wwilakehouse.dimension_city")
```

```
sale_by_date_city = df_fact_sale.alias("sale") \
.join(df_dimension_date.alias("date"), df_fact_sale.InvoiceDateKey == df_dimension_date.Date, "inner") \
.join(df_dimension_city.alias("city"), df_fact_sale.CityKey == df_dimension_city.CityKey, "inner") \
.select("date.Date", "date.CalendarMonthLabel", "date.Day", "date.ShortMonth", "date.CalendarYear", "city.City", "city.StateProvince",
"city.SalesTerritory", "sale.TotalExcludingTax", "sale.TaxAmount", "sale.TotalIncludingTax", "sale.Profit") \
.groupBy("date.Date", "date.CalendarMonthLabel", "date.Day", "date.ShortMonth", "date.CalendarYear", "city.City", "city.StateProvince",
"city.SalesTerritory") \
.sum("sale.TotalExcludingTax", "sale.TaxAmount", "sale.TotalIncludingTax", "sale.Profit") \
```

```

.withColumnRenamed("sum(TotalExcludingTax)", "SumOfTotalExcludingTax")\
.withColumnRenamed("sum(TaxAmount)", "SumOfTaxAmount")\
.withColumnRenamed("sum(TotalIncludingTax)", "SumOfTotalIncludingTax")\
.withColumnRenamed("sum(Profit)", "SumOfProfit")\
.orderBy("date.Date", "city.StateProvince", "city.City")

```

sale_by_date_city.write.mode("overwrite").format("delta").option("overwriteSchema", "true").save("Tables/aggregate_sale_by_date_city")

21. **Approach #2** – Use Spark SQL to join and aggregates data for generating business aggregates. In the below code, you are creating a temporary Spark view by joining 3 tables, doing group by to generate aggregation, renaming few of the columns. Finally, you are reading from the temporary Spark view and finally writing it as delta table in the Tables section of the lakehouse to persist with the data.

Approach #2 - sale_by_date_employee

In this cell, you are creating a temporary Spark view by joining 3 tables, doing group by to generate aggregation, renaming few of the columns.

```

1  %%sql
2  CREATE OR REPLACE TEMPORARY VIEW sale_by_date_employee
3  AS
4  SELECT
5      DD.Date, DD.CalendarMonthLabel
6      , DD.Day, DD.ShortMonth Month, CalendarYear Year
7      , DE.PreferredName, DE.Employee
8      , SUM(FS.TotalExcludingTax) SumOfTotalExcludingTax
9      , SUM(FS.TaxAmount) SumOfTaxAmount
10     , SUM(FS.TotalIncludingTax) SumOfTotalIncludingTax
11     , SUM(Profit) SumOfProfit
12 FROM wwilakehouse.fact_sale FS
13 INNER JOIN wwilakehouse.dimension_date DD ON FS.InvoiceDateKey = DD.Date
14 INNER JOIN wwilakehouse.dimension_Employee DE ON FS.SalespersonKey = DE.EmployeeKey
15 GROUP BY DD.Date, DD.CalendarMonthLabel, DD.Day, DD.ShortMonth, DD.CalendarYear, DE.PreferredName, DE.Employee
16 ORDER BY DD.Date ASC, DE.PreferredName ASC, DE.Employee ASC

```

✓ 10 sec - Command executed in 10 sec 136 ms by Arshad Ali on 4:16:26 PM, 4/14/23

Spark SQL

Spark jobs (3 of 3 succeeded)

No data available

+ Code + Markdown

In this cell, you are reading from the temporary Spark view created in the previous cell and and finally writing it as delta table in the Tables section of the lakehouse.

```

1  sale_by_date_employee = spark.sql("SELECT * FROM sale_by_date_employee")
2  sale_by_date_employee.write.mode("overwrite").format("delta").option("overwriteSchema", "true").save("Tables/aggregate_sale_by_date_em

```

✓ 19 sec - Command executed in 19 sec 74 ms by Arshad Ali on 4:16:46 PM, 4/14/23

PySpark (Python)

```

%%sql
CREATE OR REPLACE TEMPORARY VIEW sale_by_date_employee
AS
SELECT
    DD.Date, DD.CalendarMonthLabel
    , DD.Day, DD.ShortMonth Month, CalendarYear Year
    , DE.PreferredName, DE.Employee
    , SUM(FS.TotalExcludingTax) SumOfTotalExcludingTax
    , SUM(FS.TaxAmount) SumOfTaxAmount
    , SUM(FS.TotalIncludingTax) SumOfTotalIncludingTax
    , SUM(Profit) SumOfProfit
FROM wwilakehouse.fact_sale FS
INNER JOIN wwilakehouse.dimension_date DD ON FS.InvoiceDateKey = DD.Date
INNER JOIN wwilakehouse.dimension_Employee DE ON FS.SalespersonKey = DE.EmployeeKey
GROUP BY DD.Date, DD.CalendarMonthLabel, DD.Day, DD.ShortMonth, DD.CalendarYear, DE.PreferredName, DE.Employee
ORDER BY DD.Date ASC, DE.PreferredName ASC, DE.Employee ASC

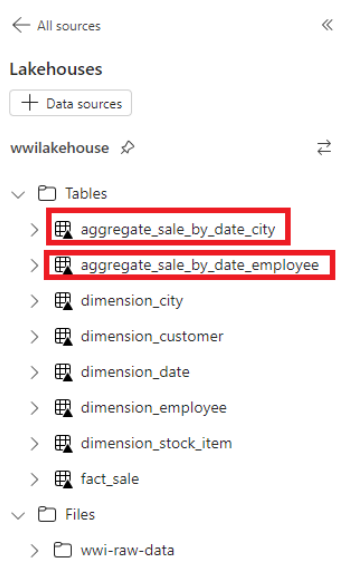
```

```

sale_by_date_employee = spark.sql("SELECT * FROM sale_by_date_employee")
sale_by_date_employee.write.mode("overwrite").format("delta").option("overwriteSchema",
"true").save("Tables/aggregate_sale_by_date_employee")

```


22. You can validate created tables by right clicking and selecting refresh on **wwilakehouse** lakehouse and you will notice these aggregate tables appear.



If you notice both these above approaches (1 and 2) produce a similar outcome, however based on developer background and preference, one approach can be picked up over the other without any need for the developer to learn a new technology and compromising on the performance.

Further, you would have noticed that you are writing data as delta lake files and then the automatic table discovery and registration feature of Fabric picks it up and registers it in the metastore. That means, you don't need to explicitly call CREATE TABLE statement to create tables to use with SQL.

Building a report

Power BI is natively integrated in the whole Fabric experience and this native integration brings unique mode of accessing the data, called DirectLake, from the lakehouse to provide the most performant query and reporting experience. DirectLake mode is a groundbreaking new engine capability to analyze very large datasets in Power BI. The technology is based on the idea of loading parquet-formatted files directly from a data lake without having to query a data warehouse or lakehouse endpoint and without having to import or duplicate data into a Power BI dataset. DirectLake is a fast path to load the data from the data lake straight into the Power BI engine, ready for analysis.

In traditional DirectQuery mode, the Power BI engine queries the data directly from the data source every time it is queried and hence query performance depends on the speed data can be retrieved from the data source. It avoids having to copy the data i.e., any changes at the source are immediately reflected in the query results while in the import mode, on the other hand, performance is much better because the data is readily available in memory without having to query the data source each time, but the Power BI engine must first copy the data into the dataset at refresh time. Any changes at the source are only picked up during the next data refresh.

DirectLake mode now eliminates this import requirement by loading the data files directly into memory. Because there is no explicit import process, it is possible to pick up any changes at the source as they occur, thus combining the advantages of DirectQuery and import mode while avoiding their disadvantages. DirectLake mode is therefore the ideal choice for analyzing very large datasets as well as datasets with frequent updates at the source.

1. Please go the **wwilakehouse** lakehouse, click **SQL analytics endpoint** under **Lakehouse** on top right of the screen to open warehouse mode of the selected lakehouse.

The screenshot shows the Microsoft Fabric Lakehouse interface. On the left, the 'Explorer' pane shows the 'wwilakehouse' lakehouse with a list of tables and files. The 'aggregate_sale_by_date_city' table is selected. On the right, the table data is displayed. A red box highlights the 'SQL analytics endpoint' button in the top right corner of the interface.

ID	Date	ARC CalendarMo...	ARC Day	ARC ShortMonth	123 CalendarYear	ARC City	ARC StateProvince	ARC SalesTerritory	SumOfTotalExclu...	SumOfTotalInclu...	SumOfTotalInclu...
1	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Bazemore	Alabama	Southeast	82944.00	12441.60	95385.60
2	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Belgreen	Alabama	Southeast	601380.00	90192.00	691472.00
3	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Coker	Alabama	Southeast	257984.00	38697.60	296681.60
4	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Eulaton	Alabama	Southeast	507584.00	76137.60	583721.60
5	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Highland Home	Alabama	Southeast	18480.00	2772.00	21252.00
6	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Jemison	Alabama	Southeast	278816.00	41822.72	320638.72
7	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Marion Junction	Alabama	Southeast	212992.00	31948.80	244940.80
8	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Nanafalia	Alabama	Southeast	312256.00	46638.40	359094.40
9	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Robertsdale	Alabama	Southeast	799222.40	119883.36	919105.76
10	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Saks	Alabama	Southeast	469312.00	70396.64	539708.64
11	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Southside	Alabama	Southeast	20160.00	3024.00	23184.00
12	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Tuscaloosa	Alabama	Southeast	36000.00	5400.00	41400.00
13	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Akhick	Alaska	Far West	601177.60	90176.64	691354.24
14	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Elvick	Alaska	Far West	106000.00	15900.00	121900.00
15	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Haycock	Alaska	Far West	284544.00	42681.60	327225.60
16	1/1/2000 12:00...	CY2000-Jan	1	Jan	2000	Ikatan	Alaska	Far West	114400.00	17160.32	131560.32

- Once you are in warehouse mode you should be able to see all the tables you created. If you don't see them yet, please click on the **Refresh** icon at the top. Next, click on **Model** tab at the bottom to open the default Power BI dataset.

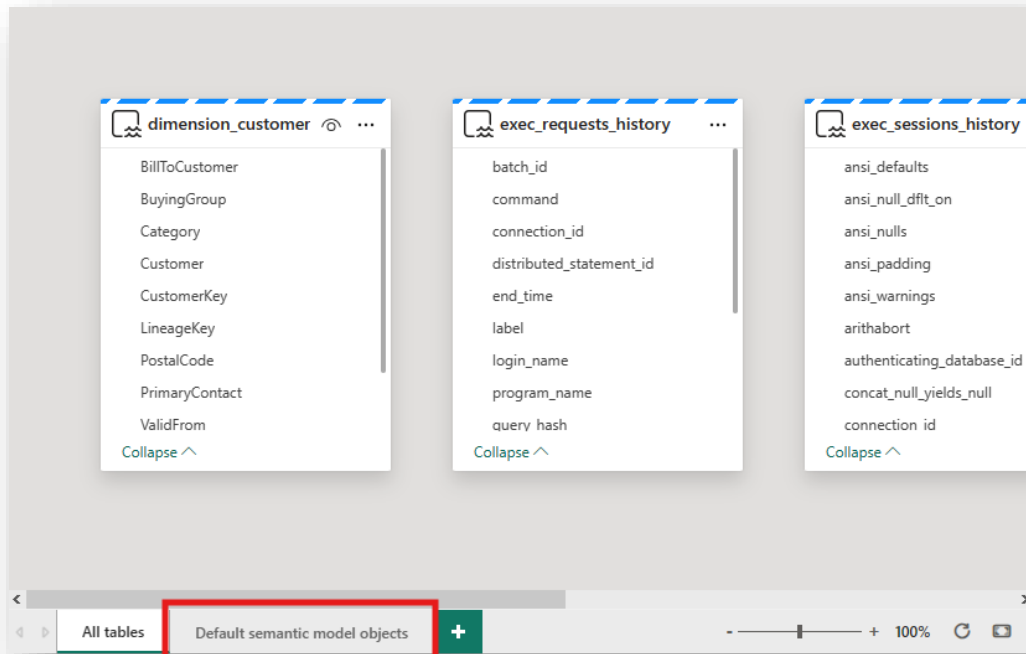
The screenshot displays the Microsoft Fabric Lakehouse interface. The top navigation bar includes 'Home', 'Reporting', and a 'SQL analytics endpoint' dropdown. Below this, a toolbar contains icons for 'New SQL query', 'New visual query', 'New report', and 'New measure'. A message states: 'This lakehouse automatically adds new objects to its default Power BI semantic model. Don't want to sync? Go to settings to turn this off. [Learn more](#)'. A 'Manage default Power BI semantic model' link is also present.

The main area is divided into three sections:


- Explorer:** A sidebar on the left showing the hierarchy of the lakehouse. It includes 'Warehouses' (wwilakehouse), 'Schemas' (dbo), 'Tables' (aggregate_sa...), 'Views', 'Functions', 'Stored Procedures', 'guest', 'INFORMATION_SCHE...', 'queryinsights', 'sys', 'Security', 'Queries', and 'My queries'. The 'Tables' section is expanded, showing 'aggregate_sa...'.
- Data preview:** A table showing 1,000 rows of data. The columns are: Date, ABC CalendarMonthLabel, ABC Day, ABC ShortMonth, 123 CalendarYear, ABC City, and ABC StateProvince. The data is filtered for '2000-01-01 00:00:00.000000'.
- Model:** A tab at the bottom of the interface, highlighted with a red box, indicating the default Power BI dataset.

The bottom status bar shows 'Succeeded (7 sec 792 ms)' and 'Columns: 12 Rows: 1,000'.

3. You need to add the newly created tables to the default semantic model. Click on the **Default semantic model objects** at the bottom of the page.

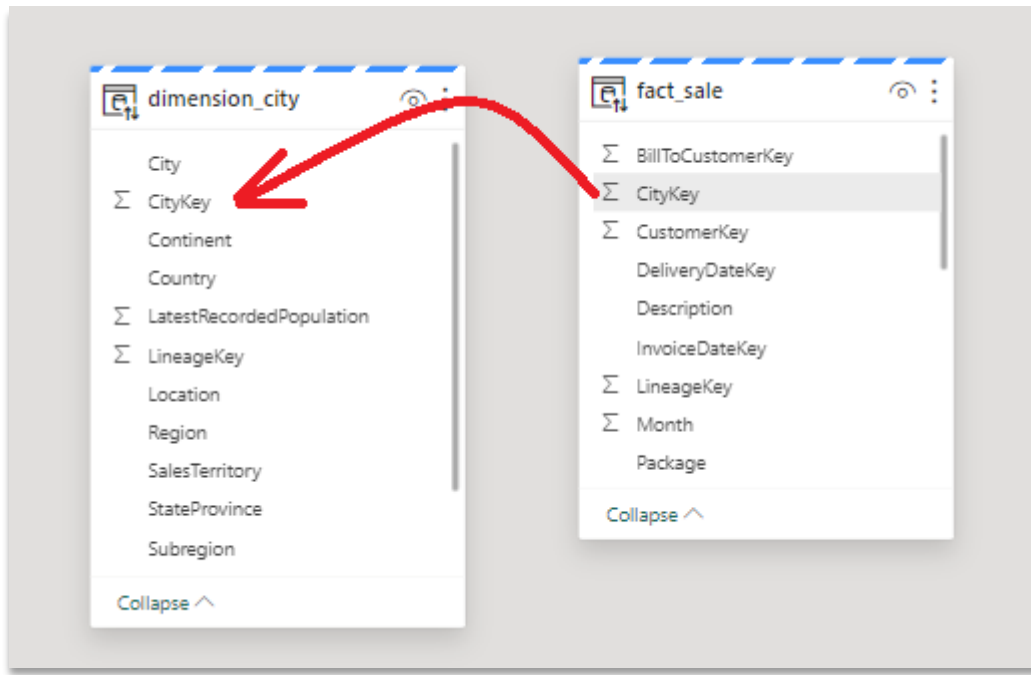


4. At the top of the page, go to the **Reporting** section, and click on **Automatically update semantic model**. This will automatically add any new tables created in this lakehouse to the default semantic model going forward. Allow a few moments for the new tables to be added to the semantic model.



Note: You can also select the **Manage default semantic model** option to specify which tables will be part of the default semantic model (but you don't need to do this now).

5. For this data model, you need to define the relationship between different tables so that you can create reports and visualizations based on data coming across different tables. From the **fact_sale** table, drag the **CityKey** field and drop it on the **CityKey** field in the **dimension_city** table to create a relationship. Check the "Assume referential integrity" and click on **Confirm** to establish the relationship. **Note** – When defining relationships, please make sure you have many to one relationship from the fact to the dimension and not vice versa.



6. On the **New Relationship** settings:
 - a. Table 1 will be populated with fact_sale and the column of CityKey.

- b. Table 2 will be populated with dimension_city and the column of CityKey.
- c. Cardinality: **Many to one (*:1)**
- d. Cross filter direction: **Single**
- e. Leave the box next to **Make this relationship active** checked.
- f. Check the box next to **Assume referential integrity**.
- g. Select **Confirm**.

New relationship [X]

Select tables and columns that are related.

Table 1	Table 2
fact_sale	dimension_city
Column: CityKey	Column: CityKey

Define cardinality and cross filter direction for tables and columns

Cardinality	Cross-filter direction
Many to one (*:1)	Single

☒ Make this relationship active

☒ Assume referential integrity. [Learn more](#)

Ok Cancel

- h. Similarly, you need to add these relationships as well:
 - StockItemKey(fact_sale) – StockItemKey(dimension_stock_item)
 - SalespersonKey(fact_sale) – EmployeeKey(dimension_employee)
 - CustomerKey(fact_sale) – CustomerKey(dimension_customer)
 - InvoiceDateKey(fact_sale) – Date(dimension_date)

Home Reporting Table tools SQL analytics endpoint

New SQL query New visual query **New report** New measure

This lakehouse automatically adds new objects to its default Power BI semantic model. Don't want to sync? Go to settings to turn this off. [Learn more](#) Manage default Power BI semantic model

Explorer

- Warehouses
- wwilakehouse
 - Schemas
 - dbo
 - Tables
 - aggregate_sale_by_date_city
 - aggregate_sale_by_date_employee
 - dimension_city
 - dimension_customer
 - dimension_date
 - dimension_employee
 - dimension_stock_item
 - fact_sale
 - SaleKey
 - CityKey
 - CustomerKey
 - BillToCustomerKey
 - StockItemKey
 - InvoiceDateKey
 - DeliveryDateKey
 - SalespersonKey
 - WWInvoiceID

fact_sale

- BillToCustomerKey
- CityKey
- CustomerKey
- DeliveryDateKey
- Description
- InvoiceDateKey
- LineageKey
- Month
- Package

dimension_date

- CalendarMonthLabel
- CalendarMonthNumber
- CalendarYearLabel
- CalendarYear
- Date
- Day
- DayNumber
- FiscalMonthLabel
- FiscalMonthNumber

dimension_city

- City
- CityKey
- Continent
- Country
- LatestRecordedPopulation
- LineageKey
- Location
- Region
- SalesTerritory
- StateProvince
- SubRegion

dimension_customer

- BillToCustomer
- BuyingGroup
- Category
- Customer
- CustomerKey
- LineageKey
- PostalCode
- PrimaryContact
- ValueFrom
- ValueTo
- WWCustomerID

dimension_employee

- Employee
- EmployeeKey
- IsSalesperson
- LineageKey
- PreferredName
- ValueFrom
- ValueTo
- WWEmployeeID

dimension_stock_item

- Barcode
- Brand
- BuyingPackage
- Color
- IsChillerStock
- LeadTimeDays
- LineageKey
- QuantityPerOuter
- RecommendedRetailPrice

aggregate sale by da...

- CalendarMonthLabel
- Date
- Day
- Employee
- Month
- PreferredName
- SumOfProfit
- SumOfTaxAmount
- SumOfTotalExcludingTax

aggregate sale by da...

- CalendarMonthLabel
- CalendarYear
- City
- Date
- Day
- SalesTerritory
- ShortMonth
- StateProvince
- SumOfProfit

exec_requests_history

- batch_id
- command
- connection_id
- distributed_statement_id
- end_time
- login_name
- program_name
- query_hash
- root_batch_id

frequently run queries

- avg_total_elapsed_time_ms
- last_statement_id
- last_run_command
- last_run_session_id
- last_run_start_time
- last_run_total_elapsed_time_ms
- max_run_total_elapsed_time_ms
- min_run_total_elapsed_time_ms
- number_of_cancelled_runs

long running queries

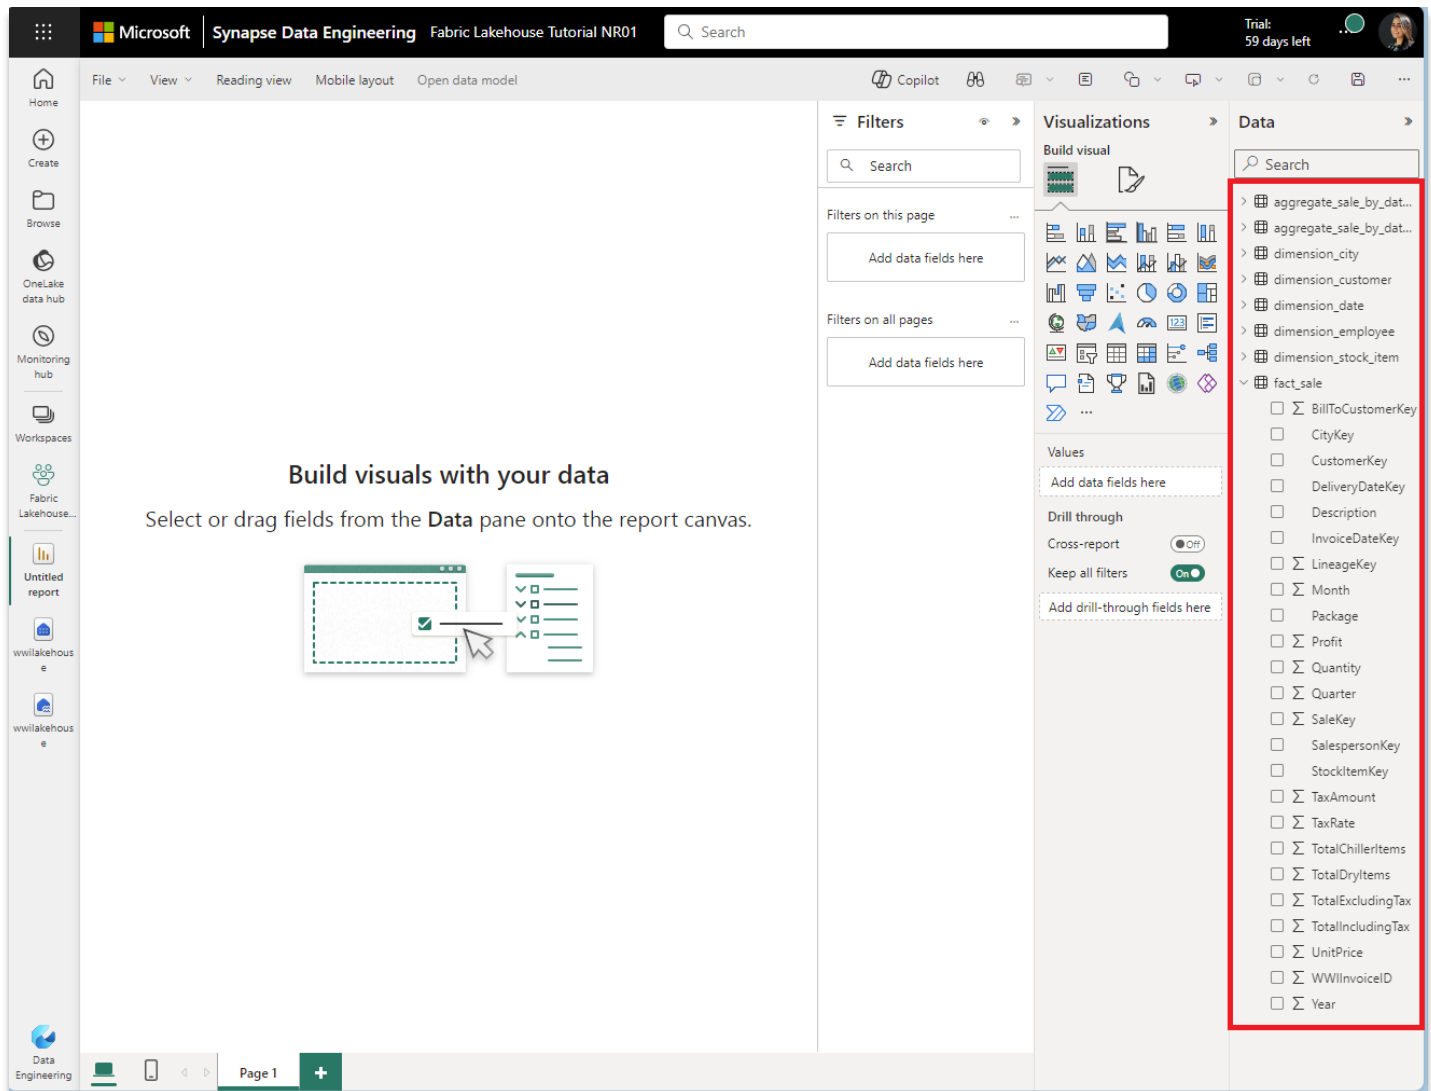
- last_statement_id
- last_run_command
- last_run_session_id
- last_run_start_time
- last_run_total_elapsed_time_ms
- median_total_elapsed_time_ms
- number_of_runs

All tables Default dataset objects + 57%

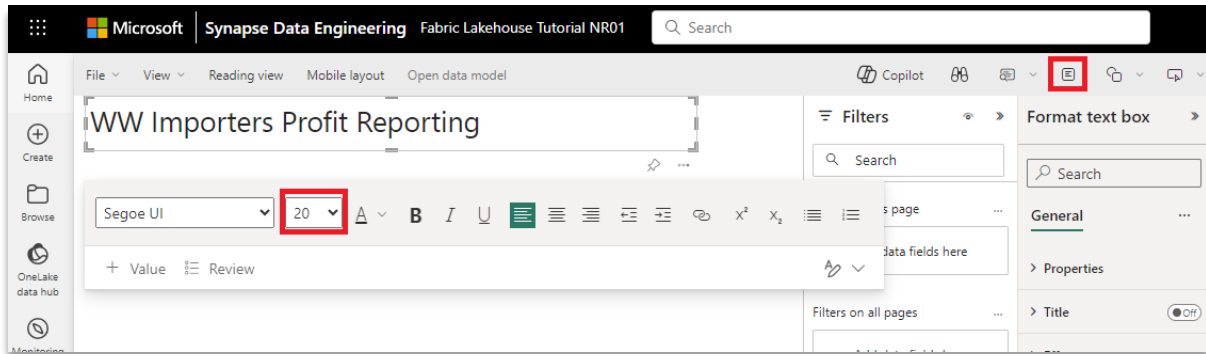
Data Query Model

After adding these above relationships, your data model is ready, as above, for reporting. Click on **New report** to start creating reports/dashboards in Power BI.

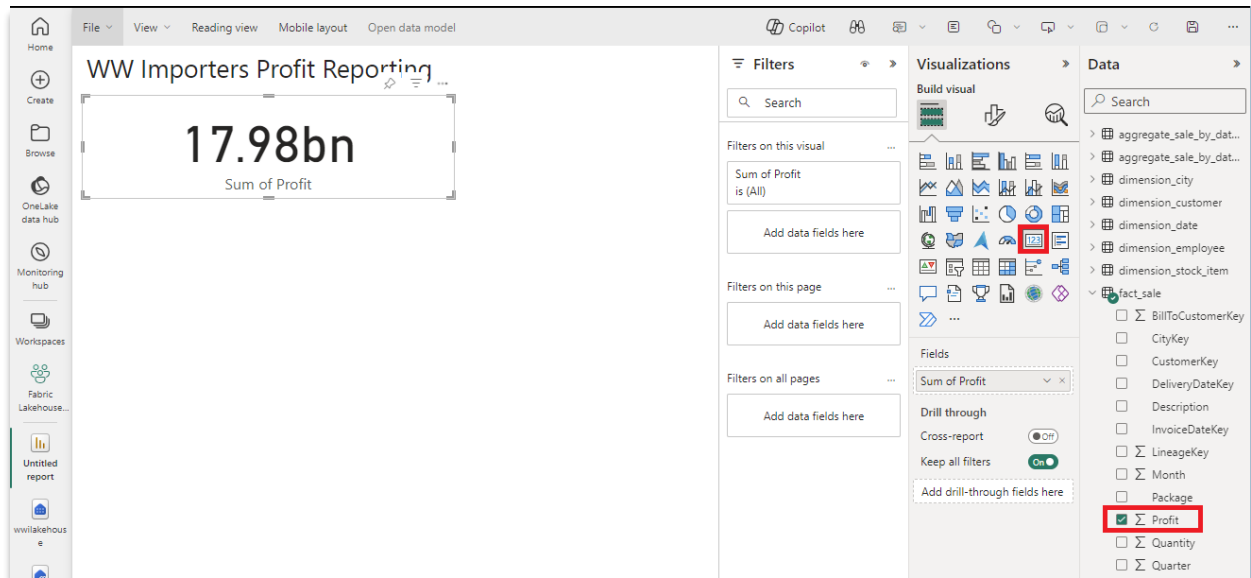
7. On the Power BI report canvas, you can create reports to meet your business requirements by dragging required columns from the **Data** pane to the canvas and using one or more of available visualizations.



8. Add a title:
 - a. In the Ribbon, click the Text Box icon
 - b. Type in **WW Importers Profit Reporting**
 - c. Highlight the text and increase size to 20 and place in the upper left of the report page

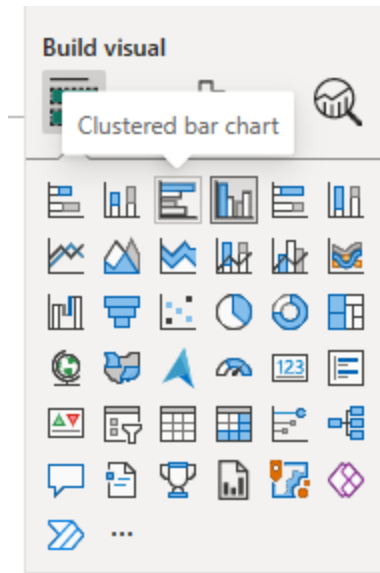


9. Add a Card
 - a. On the **Data** pane, expand **fact_sales** and check the box next to **Profit**. This will create a column chart and add the field to the Y-axis.
 - b. With the bar chart selected, click the Card visual in the visualization pane. This will convert the visual to a card.
 - c. Place the card under the title

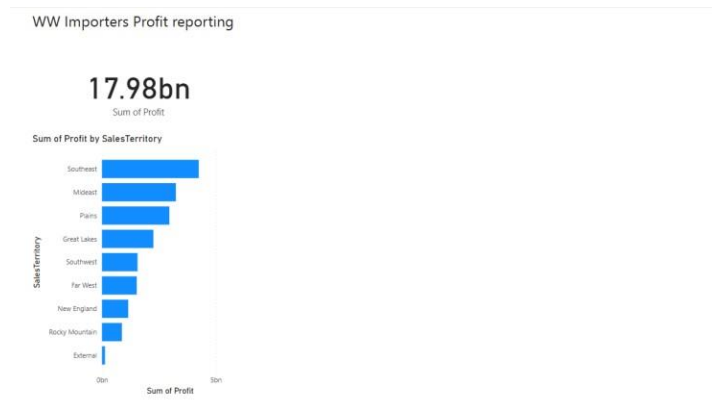


10. Add a Bar chart

- On the **Data** pane, expand **fact_sales** and check the box next to **Profit**. This will create a column chart and add the field to the Y-axis.
- On the **Data** pane, expand **dimension_city** and check the box for **SalesTerritory**. This will add the field to the X-axis.
- With the bar chart selected, click on the Clustered Bar Chart visual in the visualization pane. This will convert the column chart into a bar chart.




- Resize the Bar chart to fill in the area under the title and Card



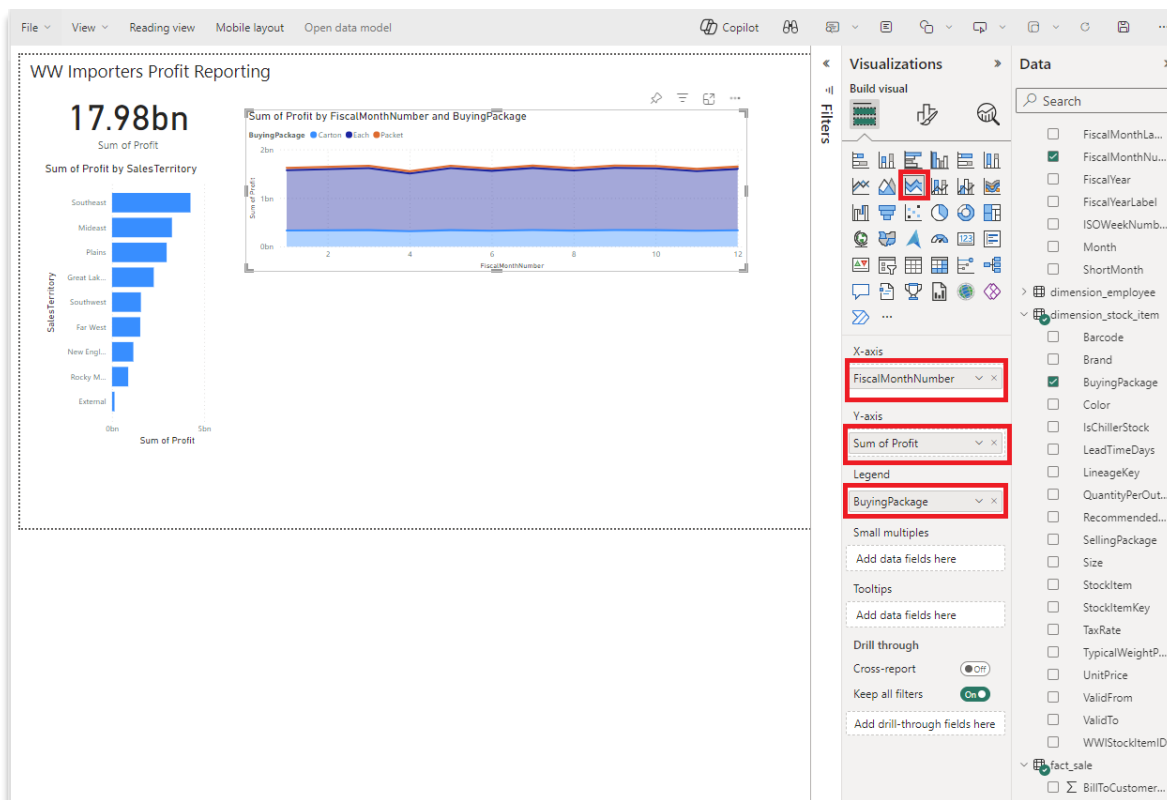
11. Click anywhere on the blank canvas (or press the Esc key) so the bar chart visual is no longer selected.

12. Build a stacked area chart visual:

- a. On the **Visualizations** pane, select the **Stacked area chart** visual.



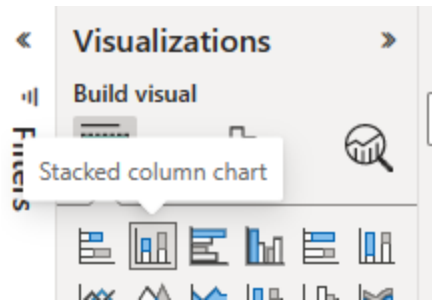
- b. Reposition and resize the stacked area chart to the right of the card and bar chart visuals created in the previous steps.
- c. On the **Data** pane, expand **fact_sales** and check the box next to **Profit**. Expand dimension_date and check the box next to **FiscalMonthNumber**. This will create a filled line chart showing profit by fiscal month.
- d. On the **Data** pane, expand **dimension_stock_item** and drag **BuyingPackage** into the Legend field well. This will add a line for each of the Buying Packages.



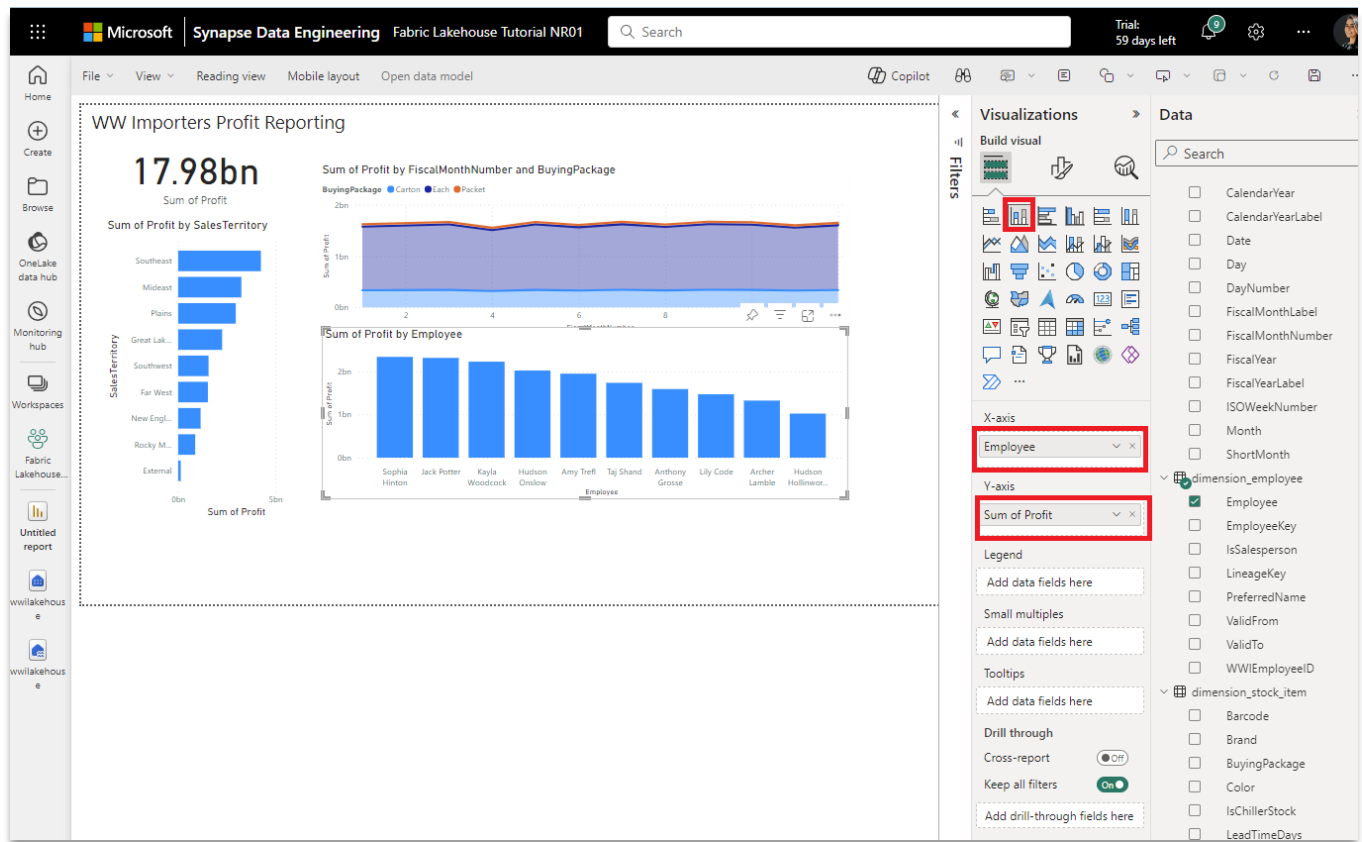
13. Click anywhere on the blank canvas (or press the Esc key) so the stacked area chart visual is no longer selected.

14. Build a column chart:

- a. On the **Visualizations** pane, select the **Stacked column chart** visual.



- b. On the **Data** pane, expand **fact_sales** and check the box next to **Profit**. This will add the field to the Y-axis.
- c. On the **Data** pane, expand **dimension_employee** and check the box next to **Employee**. This will add the field to the X-axis.



15. Click anywhere on the blank canvas (or press the Esc key) so the LINE chart visual is no longer selected.
16. From the ribbon, select **File > Save**.
17. Enter the name of your report as **Profit Reporting**.
18. Select **Save**

